



Higher-Order and Tuple-Based Massively-Parallel Prefix Sums

Sepideh Maleki*, Annie Yang, and Martin Burtscher

Department of Computer Science

TEXAS  STATE
UNIVERSITY

The rising STAR of Texas

ECL

Efficient Computing Laboratory



Prefix Sum

- Given an array of values (integer or real values)

0	1	2	3	4	5	6	7
3	2	0	7	-6	1	-9	5

- Compute the array whose elements are the sum of **all previous elements** from the original array

0	1	2	3	4	5	6	7
3	5	5	12	6	7	-2	3

- A prefix scan is a **generalization** of the prefix sum where the operation doesn't have to be addition

Some Scan Operators

Operator	Identity element	Example
+	0	$X + 0 = X$
Min	Maximum Representative value	$\text{Min}(X, \infty) = X$
Max	Minimum Representative value	$\text{Max}(X, -\infty) = X$
Multiply	1	$X * 1 = X$
Logical Or	FALSE	$X \ \ \text{FALSE} = X$
Logical AND	TRUE	$X \ \&\& \ \text{TRUE} = X$

Uses of Prefix Sums and Scans


- Fundamental **building block** of parallel algorithms
 - Can be computed efficiently **in parallel** in $\log(n)$ steps
 - Help parallelize many seemingly serial algorithms
- Examples
 - Buffer allocation
 - Radix sort
 - Quicksort
 - String comparison
 - Lexical analysis
 - Run-length encoding
 - Histograms
 - Polynomial evaluation
 - Stream compaction
 - **Data compression**

Highlights

- GPU-friendly algorithm for prefix scans called **SAM**
- Novelties and features
 - Higher-order support that is communication optimal
 - Tuple-value support with constant workload per thread
 - Carry propagation scheme with $O(1)$ auxiliary storage
 - Implemented in unified 100-statement CUDA kernel
- Results
 - Outperforms CUB by up to 2.9-fold on higher-order and by up to 2.6-fold on tuple-based prefix sums

Data Compression

- Data compression algorithms
 - **Data model** predicts next value in input sequence and emits difference between actual and predicted value
 - **Coder** maps frequently occurring values to produce shorter output than infrequent values
- Delta encoding
 - Widely used data model
 - Computes **difference sequence** (i.e., predicts current value to be the same as previous value in sequence)
 - Used in image compression, speech compression, etc.



Delta Coding

- Delta **encoding** is embarrassingly parallel
- Delta **decoding** appears to be sequential
 - Decoded prior value needed to decode current value
- **Prefix sum** decodes delta encoded values
 - Decoding can also be done in parallel

Input sequence	1, 2, 3, 4, 5, 2, 4, 6, 8, 10
Difference sequence (encoding)	1, 1, 1, 1, 1, -3, 2, 2, 2, 2
Prefix sum (decoding)	1, 2, 3, 4, 5, 2, 4, 6, 8, 10

Extensions of Delta Coding

■ Higher orders

- Higher-order predictions are often more accurate

- First order

$$\text{out}_k = \text{in}_k - \text{in}_{k-1}$$

- Second order

$$\text{out}_k = \text{in}_k - 2 \cdot \text{in}_{k-1} + \text{in}_{k-2}$$

- Third order

$$\text{out}_k = \text{in}_k - 3 \cdot \text{in}_{k-1} + 3 \cdot \text{in}_{k-2} - \text{in}_{k-3}$$

■ Tuple values

- Data frequently appear in tuples

- Two-tuples

$$x_0, y_0, x_1, y_1, x_2, y_2, \dots, x_{n-1}, y_{n-1}$$

- Three-tuples

$$x_0, y_0, z_0, x_1, y_1, z_1, \dots, x_{n-1}, y_{n-1}, z_{n-1}$$

Problem and Solution

- Conventional prefix sums are insufficient
 - **Do not** decode higher-order delta encodings
 - **Do not** decode tuple-based delta encodings
- Prior work
 - Requires **inefficient workarounds** to handle higher-order and tuple-based delta encodings
- SAM algorithm and implementation
 - Directly and **efficiently** supports these generalizations
 - Even supports **combination** of higher orders and tuples

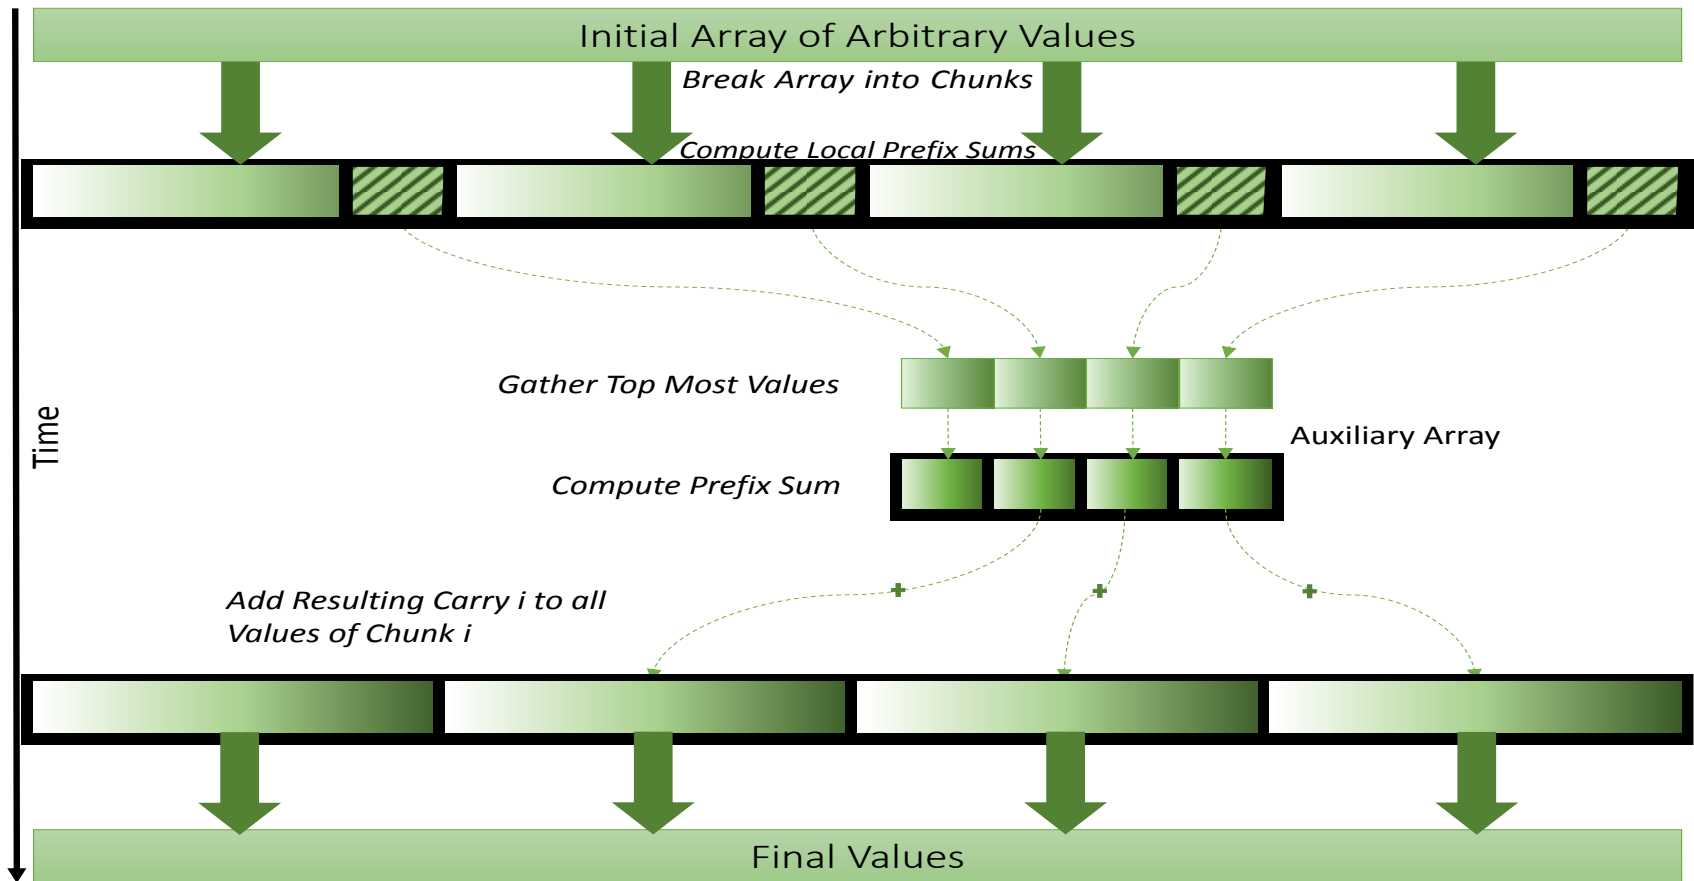
Work Efficiency of Prefix Sums

- Sequential prefix sum requires only a single pass
 - $2n$ data movement through memory
 - Linear $O(n)$ complexity

```
1 out[0] = 0
2 for i from 1 to n do
3     out[i] = out[i - 1] + in[i - 1]
```

- Parallel algorithm should have same complexity
 - $O(n)$ applications of the sum operator

Hierarchical Parallel Prefix Sum



Standard Prefix-Sum Implementation

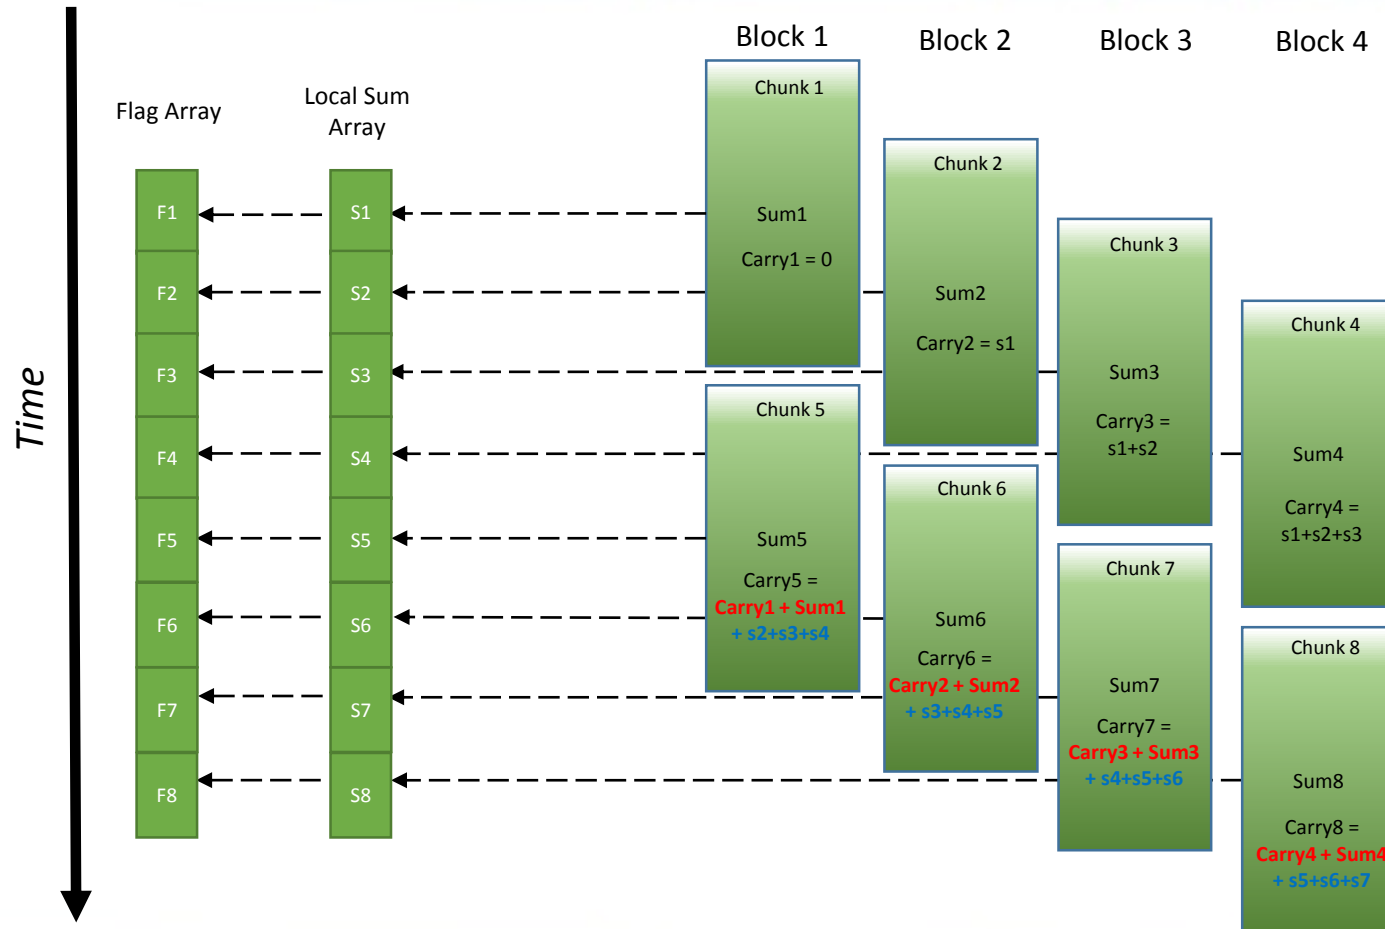
- Based on 3-phase approach
- Reads and writes every element twice
 - $4n$ main-memory accesses
- Auxiliary array is stored in global memory
 - Calculation is performed across blocks
- High-performance implementations
 - Allocate and process several values per thread
- **Thrust** and **CUDPP** use this hierarchical approach

SAM Base Implementation

- Intra-block prefix sums
 - Computes prefix sum of each chunk conventionally
 - Writes **local sum** of each chunk to auxiliary array
 - Writes **ready flag** to second auxiliary array
- Inter-block prefix sums
 - Reads local sums of **all** prior chunks
 - **Adds up local sums to calculate carry**
 - Updates all values in chunk using carry
 - Writes final result to global memory



Pipelined Processing of Chunks



Carry Propagation Scheme

- Persistent-block-based implementation
 - Same block processes every k^{th} chunk
 - Carries require only **$O(1)$ computation** per chunk
- Circular-buffer-based implementation
 - Only $3k$ elements needed at any point in time
 - Local sums and ready flags require **$O(1)$ storage**
- Redundant computations for latency hiding
 - Write-followed-by-independent-reads pattern
 - Multiple values processed per thread (fewer chunks)



Higher-order Prefix Sums

Higher-order Prefix Sums

- Higher-order difference sequences can be computed by **repeatedly** applying first order

Input values	1, 2, 3, 4, 5, 2, 4, 6, 8, 10
First order	1, 1, 1, 1, 1, -3, 2, 2, 2, 2
Second order	1, 0, 0, 0, 0, -4, 5, 0, 0, 0

- Prefix sum is the inverse of order-1 differencing
 - K prefix sums will decode an order- k sequence
- No direct solution for computing higher orders
 - Must use iterative approach
 - Other codes' memory accesses **proportional** to order

Higher-order Prefix Sums (cont.)

- SAM is more **efficient**
 - Internally iterates only the computation phase
 - Does not read and write data in each iteration
 - Requires only $2n$ main-memory accesses for any order
- SAM's higher-order implementation
 - Employs **multiple sum arrays**, one per order
 - Each sum array is an $O(1)$ circular buffer
 - Needs $O(1)$ **non-Boolean** ready 'flags'
 - Uses counts to indicate iteration of current local sum



Tuple-based Prefix Sums

Tuple-based Prefix Sums

- Data may be tuple based $x_0, y_0, x_1, y_1, \dots, x_{n-1}, y_{n-1}$
- Other codes have to handle tuples as follows
 - **Reordering** elements, compute, undo reordering
 - **Slow** due to reordering and may require extra memory

$$\begin{array}{l} x_0, x_1, \dots, x_{n-1} \mid y_0, y_1, \dots, y_{n-1} \\ \Sigma_0^0 x_i, \Sigma_0^1 x_i, \dots, \Sigma_0^{n-1} x_i \mid \Sigma_0^0 y_i, \Sigma_0^1 y_i, \dots, \Sigma_0^{n-1} y_i \\ \Sigma_0^0 x_i, \Sigma_0^0 y_i, \Sigma_0^1 x_i, \Sigma_0^1 y_i, \dots, \Sigma_0^{n-1} x_i, \Sigma_0^{n-1} y_i \end{array}$$

- Defining a tuple **data type** as well as the plus **operator**
 - **Slow** for large tuples due to excessive register pressure

Tuple-based Prefix Sums (cont.)

- SAM is more **efficient**
 - No reordering
 - No special data types or overloaded operators
 - Always same amount of data per thread
- SAM's tuple implementation
 - Employs **multiple sum arrays**, one per tuple element
 - Each sum array is an $O(1)$ circular buffer
 - Uses modulo operations to determine which array to use
 - Still employs single $O(1)$ Boolean flag array

Experimental Methodology

- Evaluate following prefix sum implementations

- **Thrust** library (from CUDA Toolkit 7.5)

- $4n$ memory accesses

- **CUDPP** library 2.2

- $4n$ memory accesses

- **CUB** library 1.5.1

- $2n$ memory accesses

- **SAM** 1.1

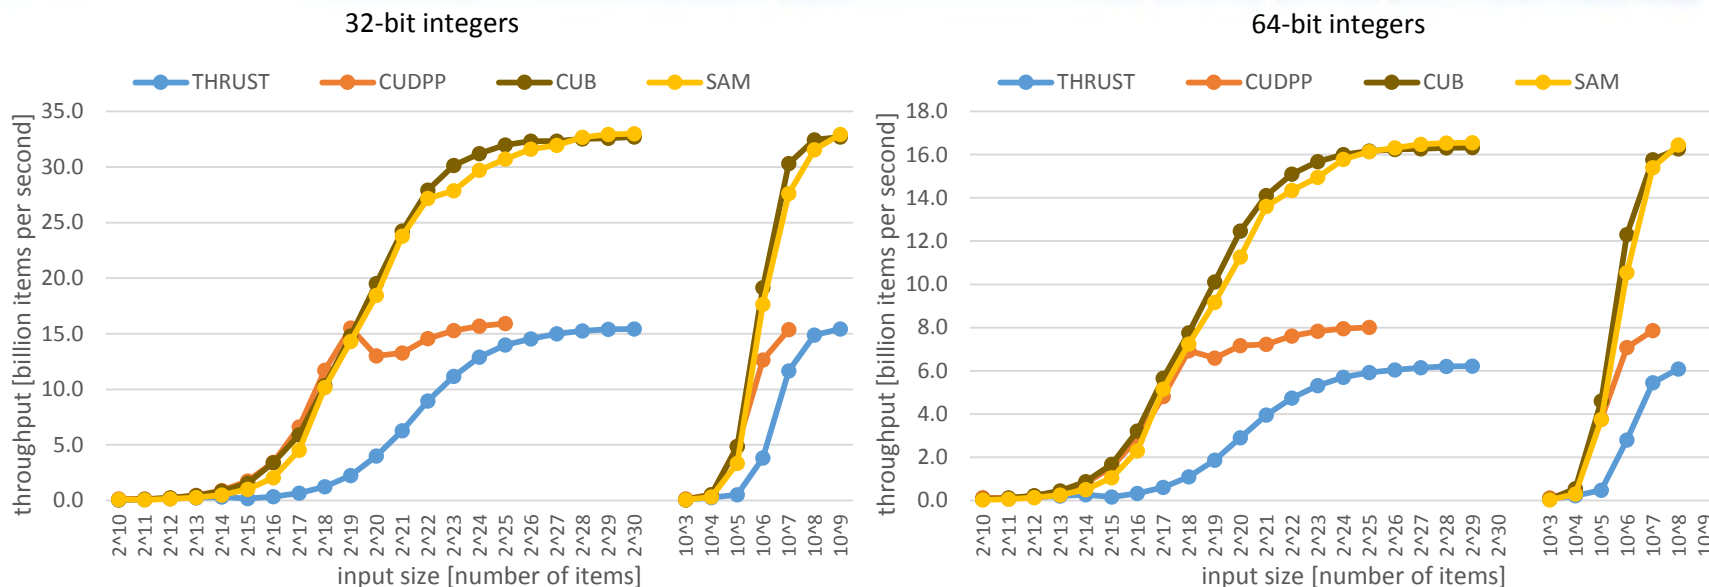
- $2n$ memory accesses

GPU	GeForce Titan X	Tesla K40c
Architecture	Maxwell	Kepler
PE	3072	2880
Multiprocessors	24	15
Persistent Blocks	48	30
Global Memory	12 GB	12 GB
Peak Bandwidth	336 GB/s	288 GB/s



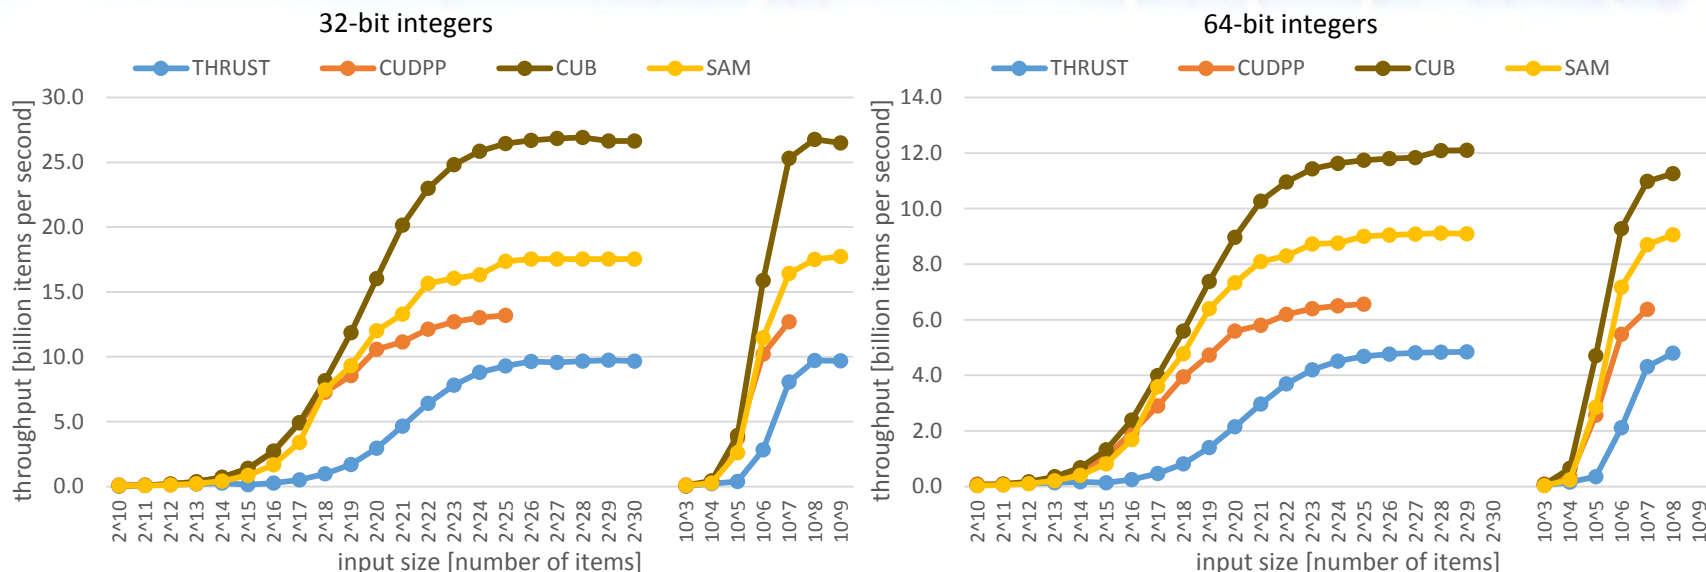
Performance Evaluation

Prefix Sum Throughputs (Titan X)




- SAM and CUB outperform the other approaches ($2n$ vs. $4n$)
- For 64-bit values, throughputs are about half (but same GB/s)
- SAM matches cudaMemcpy throughput at high end (264 GB/s)
 - Surprising since SAM was designed for higher orders and tuples

Prefix Sum Throughputs (K40)




- K40 throughputs are lower for all algorithms
- SAM is faster than Thrust/CUDPP on medium and large inputs
- CUB outperforms SAM by 50% on large inputs on 32-bits ints
 - SAM's implementation is not a particularly good fit for this older GPU

Higher-order Throughputs (Titan X)




- Throughputs decrease as order increases due to more iterations
- **SAM's performance advantage increases with higher orders**
 - Always executes $2n$ global memory accesses
 - Outperforms CUB by 52% on order 2, 78% on order 5, and 87% on order 8

Higher-order Throughputs (K40)




- CUB outperforms SAM on orders 2 and 5, but not on order 8
 - Again, **SAM's relative performance increases with higher orders**
- SAM's relative performance over CUB is higher on 64-bit values
 - Baseline advantage of CUB over SAM is smaller for 64-bit values

Tuple-based Throughputs (Titan X)



- Throughputs decrease with larger tuple sizes due to extra work
- **SAM's performance advantage increases with larger tuple sizes**
 - Larger tuples increase register pressure in CUB but **not** in SAM
 - SAM is 17% slower on 2-tuples but 20% faster on 5-tuples and 34% faster on 8-tuples

Tuple-based Throughputs (K40)



- SAM outperforms CUB on 8-tuples (and larger tuples)
 - Again, **SAM's relative performance increases with larger tuple sizes**
- The benefit of SAM over CUB is higher with 64-bit values
 - SAM already outperforms CUB on 5-tuples

Summary

- SAM directly supports generalized **prefix scans**
 - **Higher-order** prefix scans
 - **Tuple-based** prefix scans
- SAM performance on Maxwell and Kepler GPUs
 - Reaches **cudaMemcpy throughput** on large inputs
 - Outperforms all alternatives by up to 2.9x on order-eight and by up to 2.6x on eight-tuple prefix sums
 - **SAM's relative performance increases with higher orders and larger tuple sizes**

Question?

- Contact Info: Smaleki@txstate.edu



<http://cs.txstate.edu/~burtscher/research/SAM/>

- Acknowledgments
 - National Science Foundation
 - NVIDIA Corporation
 - Texas Advanced Computing Center

