

Coverage-Directed Differential Testing of JVM Implementations

Yuting Chen¹, Ting Su², Chengnian Sun³, Zhendong Su³, and Jianjun Zhao^{1,4}

¹Shanghai Jiao Tong University

²East China Normal University

³University of California, Davis

⁴Kyushu University

Outline


- Motivation
 - Testing of JVMs
 - Test redundancy
 - Goal + Key Observations
- Design
- Evaluation
- Related Work and Conclusion

Background: JVM



JVM Testing

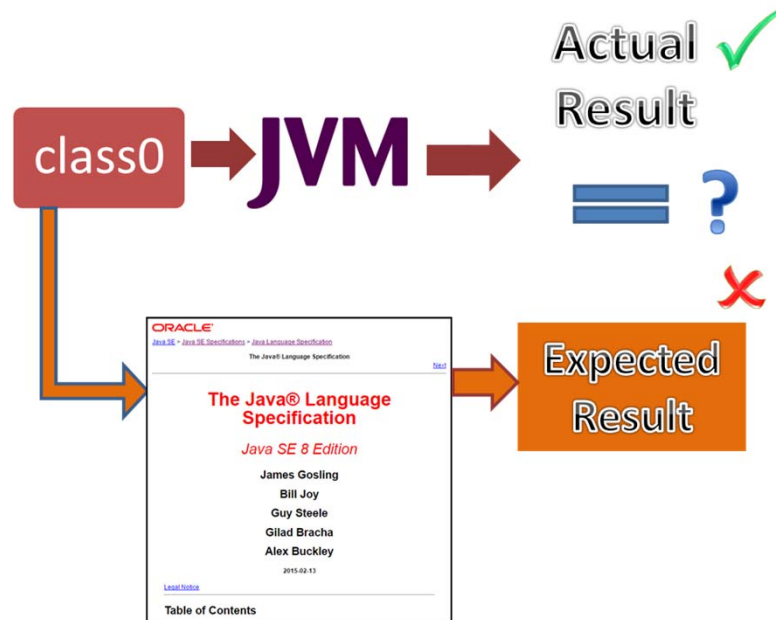
- Testing a JVM using a number of test classfiles




- ① How is a JVM defect exposed?
- ② How is a test classfile achieved?

Challenge 1: How to expose a JVM defect?

Challenge:
no test oracles



Solution:
differential JVM testing



An Example of JVM Behavior Discrepancy

```
1  ...
2  MD5 checksum 8fb69050bbcb9a83ddd90ae393368c5e
3  ...
4  class M1436188543
5  minor version: 0
6  major version: 51
7  flags: ACC_SUPER
8  Constant pool:
9  ...
10 #7 = Utf8 <clinit>
11 #8 = Utf8 ()V
12 #9 = Class #19 // java/lang/System
13 #10 = Utf8 Code
14 #11 = Utf8 main
15 ...
16
17 public abstract {};
18 flags: ACC_PUBLIC, ACC_ABSTRACT
19 public static void main(java.lang.String[]);
20 flags: ACC_PUBLIC, ACC_STATIC
21 Code:
22 stack=2, locals=1, args_size=1
23 0: getstatic #12 // Field java/lang/System.out:
   Ljava/io/PrintStream;
24 3: ldc #4 // String Completed!
25 5: invokevirtual #21 // Method java/io/
   PrintStream.println:(Ljava/lang/String;)V
26 8: return
27 }
```

public abstract {};

- HotSpot takes it as a **ordinary** method
- J9 reports a **format error**

*Cause: the JVM specification says that “other methods named <clinit> in a class file are **of no consequence**. They are not class or interface initialization methods.”*

A class method needs to be more strictly defined


Challenge 2: How to obtain test classfiles?

- Option 1: using the real-world classfiles



Some classes can reveal
compatibility issues

- Option 2: domain-aware fuzz testing



An Example of JVM Behavior Discrepancy

```
1  ...
2  MD5 checksum 8fb69050bbcb9a83ddd90ae393368c5e
3  ...
4  class M1436188543
5  minor version: 0
6  major version: 51
7  flags: ACC_SUPER
8  Constant pool:
9  ...
10 #7 = Utf8 <clinit>
11 #8 = Utf8 ()V
12 #9 = Class #19 // java/lang/System
13 #10 = Utf8 Code
14 #11 = Utf8 main
15 ...
16 {
17   public abstract {};
18   flags: ACC_PUBLIC, ACC_ABSTRACT
19   public static void main(java.lang.String[]);
20   flags: ACC_PUBLIC, ACC_STATIC
21   Code:
22     stack=2, locals=1, args_size=1
23     0: getstatic #12 // Field java/lang/System.out:
        Ljava/io/PrintStream;
24     3: ldc #4 // String Completed!
25     5: invokevirtual #21 // Method java/io/
        PrintStream.println:(Ljava/lang/String;)V
26     8: return
27 }
```

public abstract mymethod{};

clinit

More JVM discrepancies are revealed by
domain-aware fuzz testing


Challenge 2: How to obtain test classfiles?

- Option 1: using the real-world classfiles



Some classes can reveal
compatibility issues

- Option 2: domain-aware fuzz testing




Challenge 2: How to obtain test classfiles?



- An infinite number of test classfiles can be created
- They may reveal a small number of JVM discrepancies

- Solution 2: domain-aware fuzz testing




Key Observation (1)

- A classfile can encompass intricate constraints
 - Corner cases can be created through rewriting seeds

legal

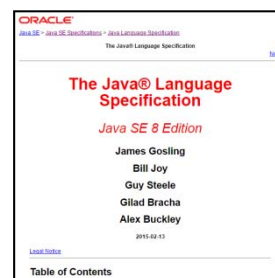


Syntax tree



Revised syntax tree


legal
or
illegal




JVM

Key Observation (2)

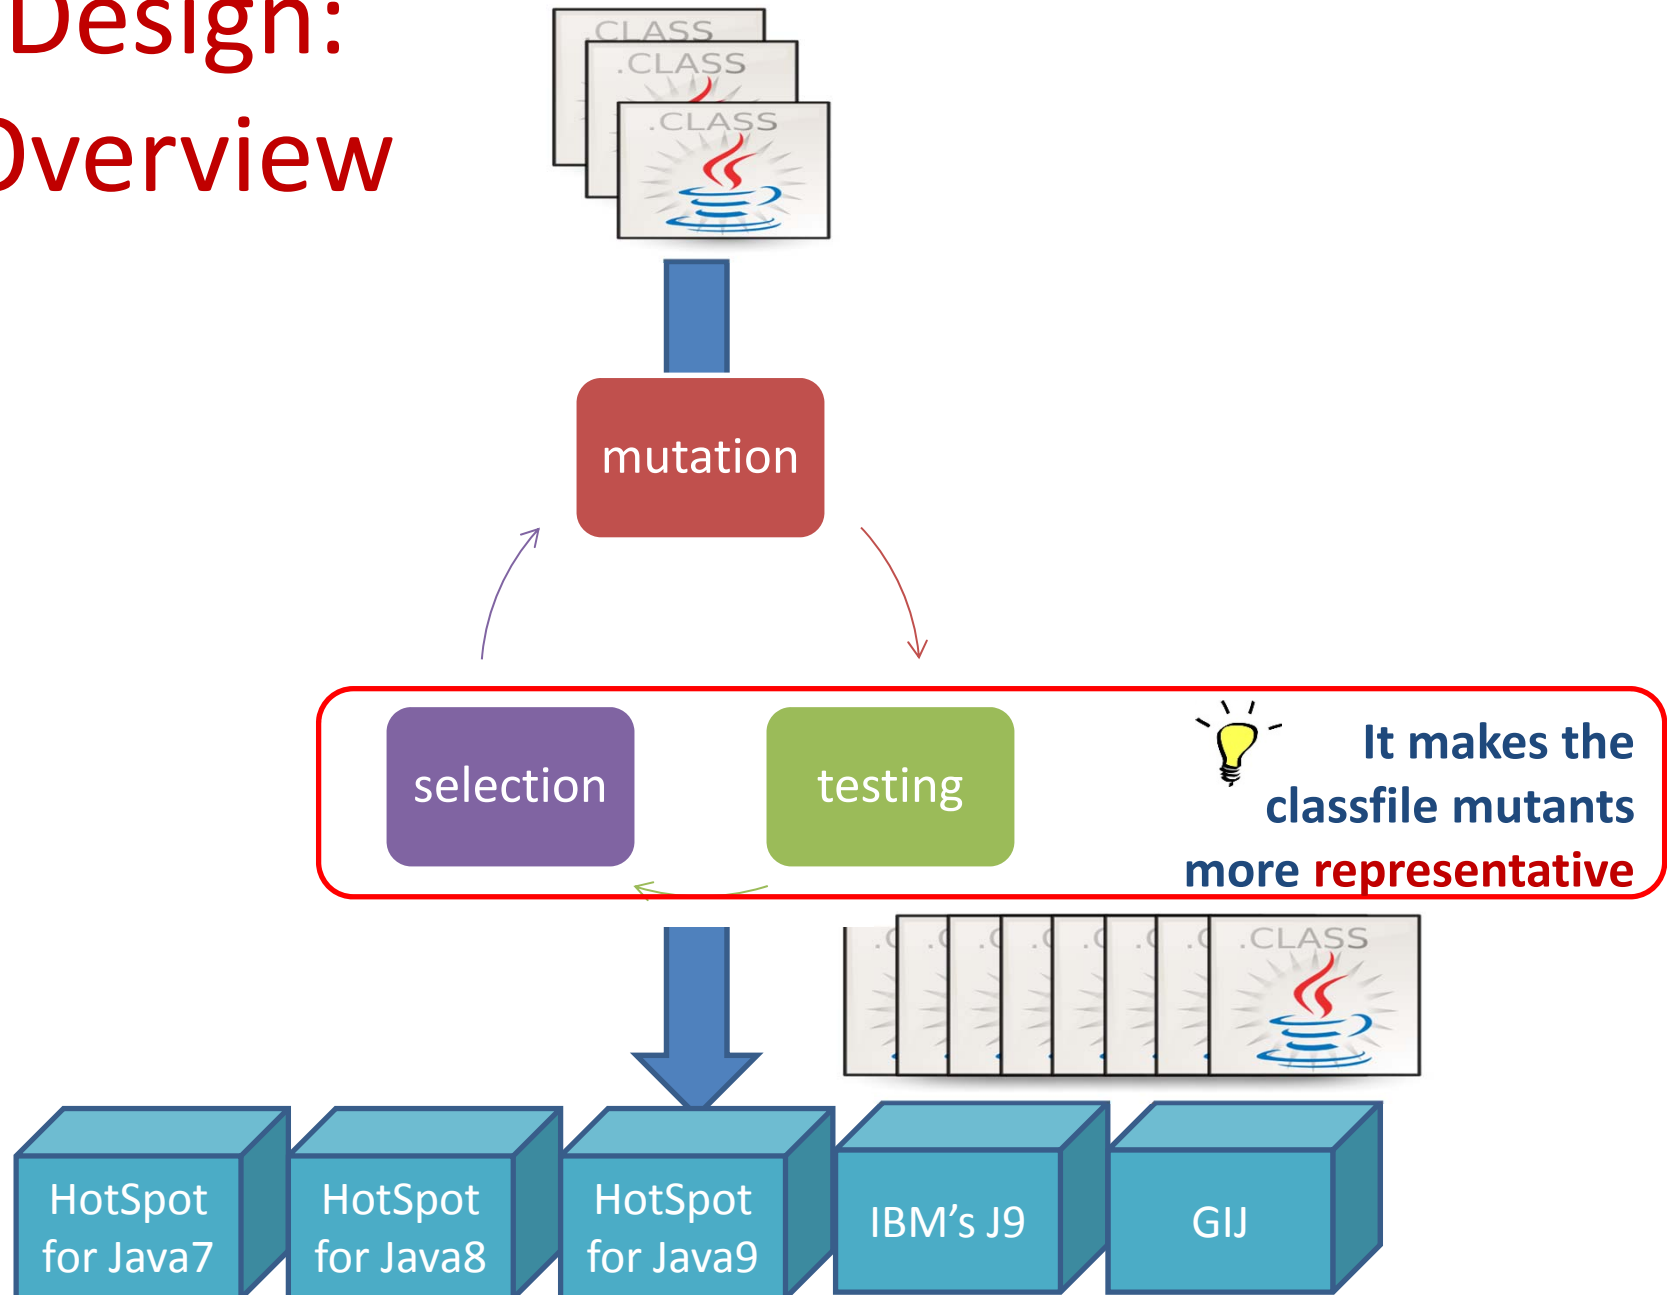
- Equivalence class partition (ECP) saves the testing cost
 - ECP works only if we can decide whether two tests belong to the same partition



Our Design: An Overview






Our Design: An Overview



Outline


- Motivation
- Design
 - Mutating classfiles
 - Selecting representative classfile mutants
 - Selectively applying mutators
 - Differentially testing JVMs
- Evaluation
- Related Work and Conclusion

Mutating Classfiles

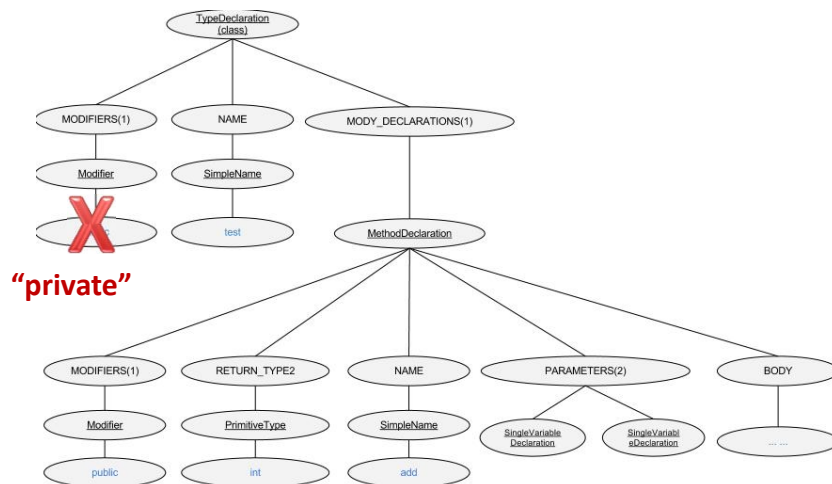


We have designed **129** mutators
for mutating classfiles

Mutating Classfiles (2)



123 mutators are designed for rewriting the ASTs of the seeds

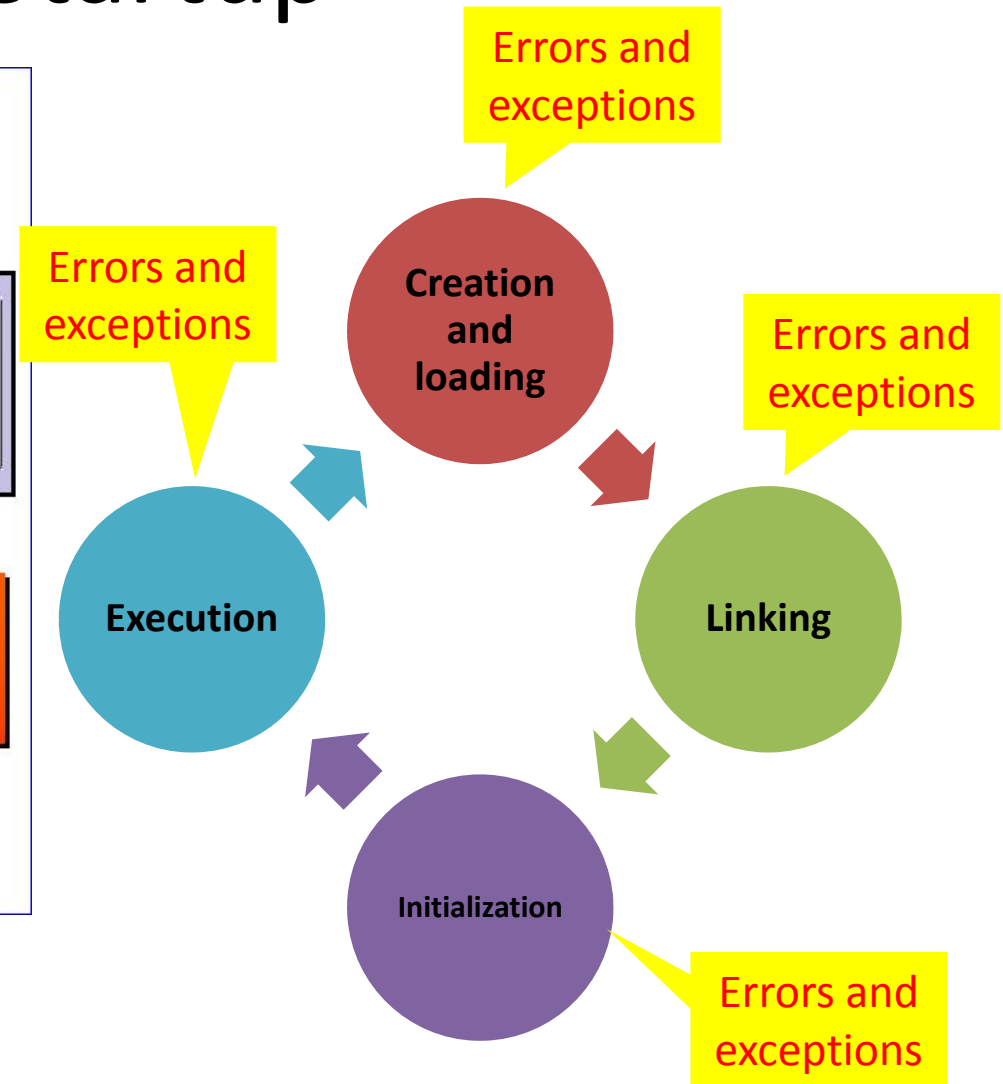
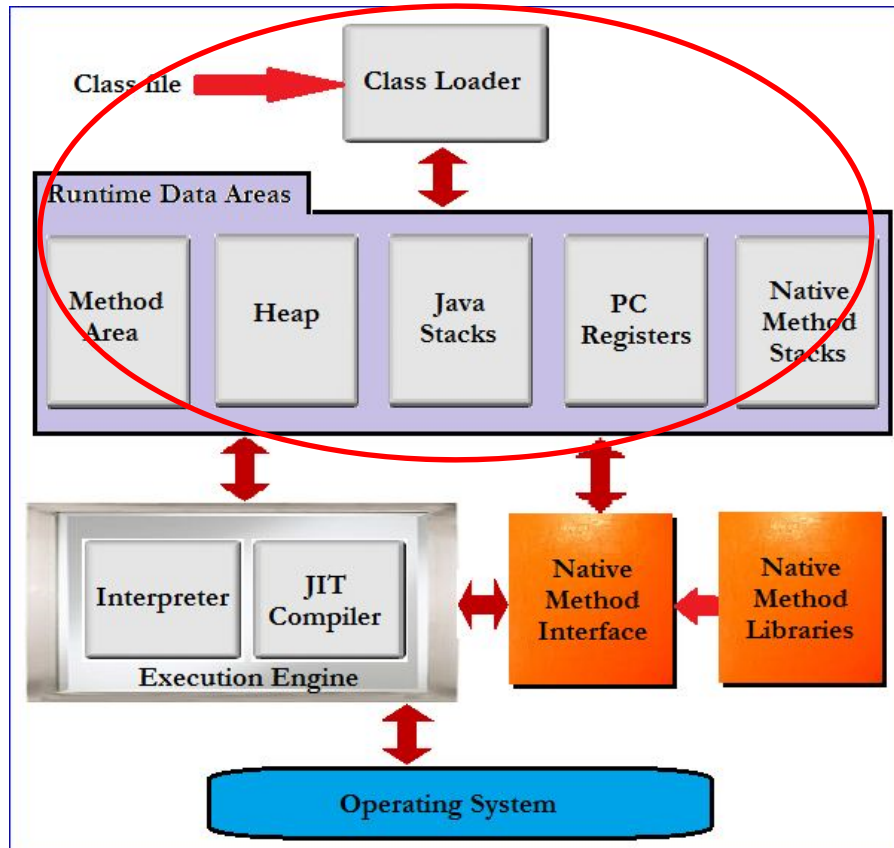


Six mutators are designed for rewriting the Jimple files of the seeds

```
r0:=parameter0: java.lang.String[];  
r1:=<java.lang.System: java.io.PrintStream out>;  
  
virtualinvoke $r1.<java.io.PrintStream:void  
println(java.lang.String)>("Executed");  
  
...
```

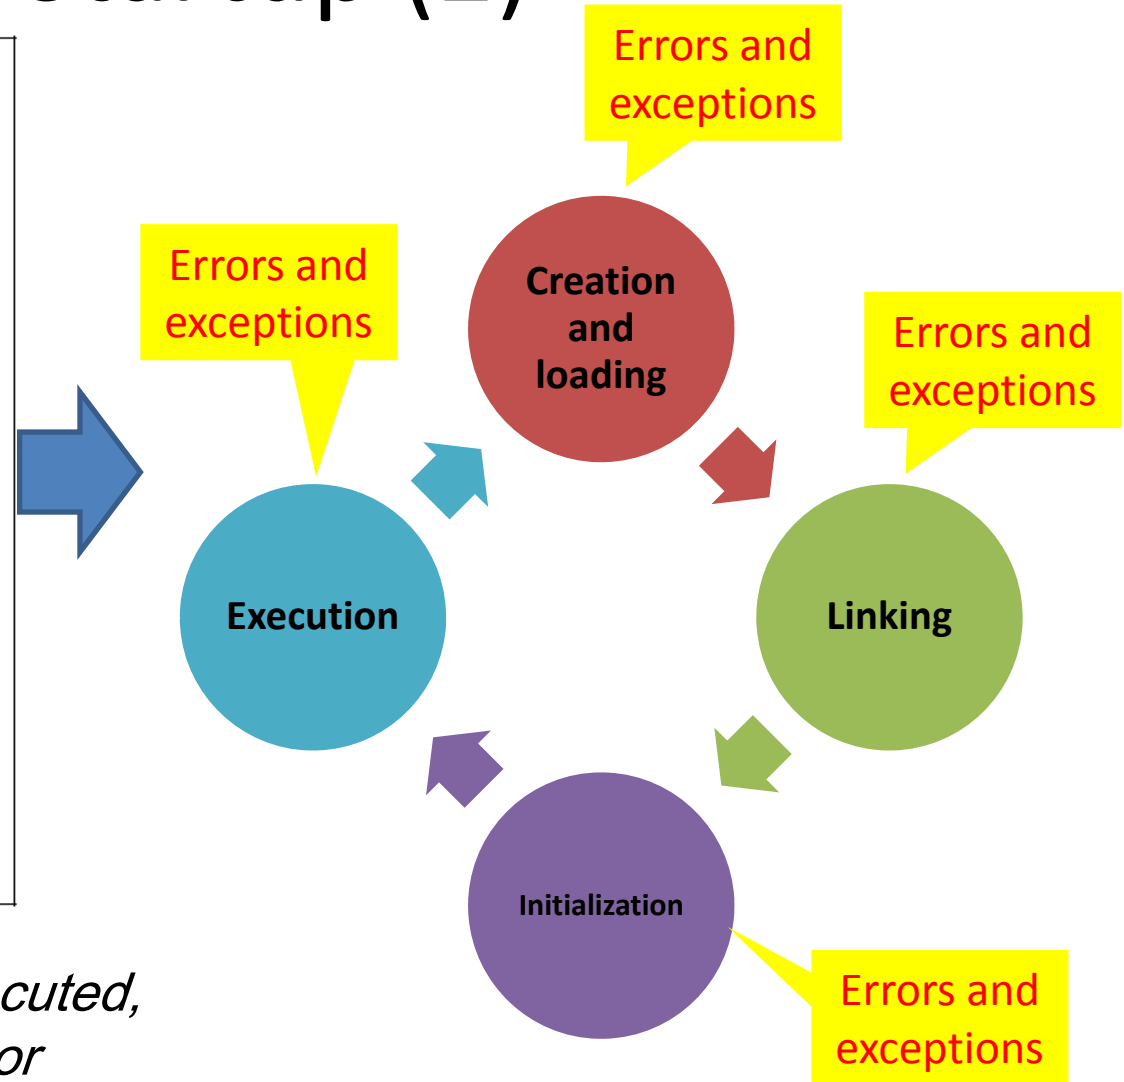
- **Limitation:** Only the JVMs' startup processes can be tested
 - The mutated program constructs/attributes may be less likely to be activated during execution

JVM Startup



JVM Startup (2)


```
...
MD5 checksum 8fb69050bbcb9a83ddd90ae393368c5e
...
class M1436188543
  minor version: 0
  major version: 51
  flags: ACC_SUPER
  Constant pool:
    ...
    #7 = Utf8 <clinit>
    #8 = Utf8 ()V
    #9 = Class #19 // java/lang/System
    #10 = Utf8 Code
    #11 = Utf8 main
    ...
  {
    public abstract {};
    flags: ACC_PUBLIC, ACC_ABSTRACT
    public static void main(java.lang.String[]);
    flags: ACC_PUBLIC, ACC_STATIC
    Code:
      stack=2, locals=1, args_size=1
      0: getstatic #12 // Field java/lang/System.out:
        Ljava/io/PrintStream;
      3: ldc #4 // String Completed!
      5: invokevirtual #21 // Method java/io/
        PrintStream.println:(Ljava/lang/String;)V
      8: return
  }
}
```



*Can this class be **normally** executed, or at which stage some **errors** or **exceptions** can be thrown out?*


Outline

- Motivation
- Design
 - Mutating classfiles
 - Selecting representative classfile mutants
 - Selectively applying mutators
 - Differentially testing JVMs
- Evaluation
- Related Work and Conclusion





- ECP

Selecting Representative Classfiles



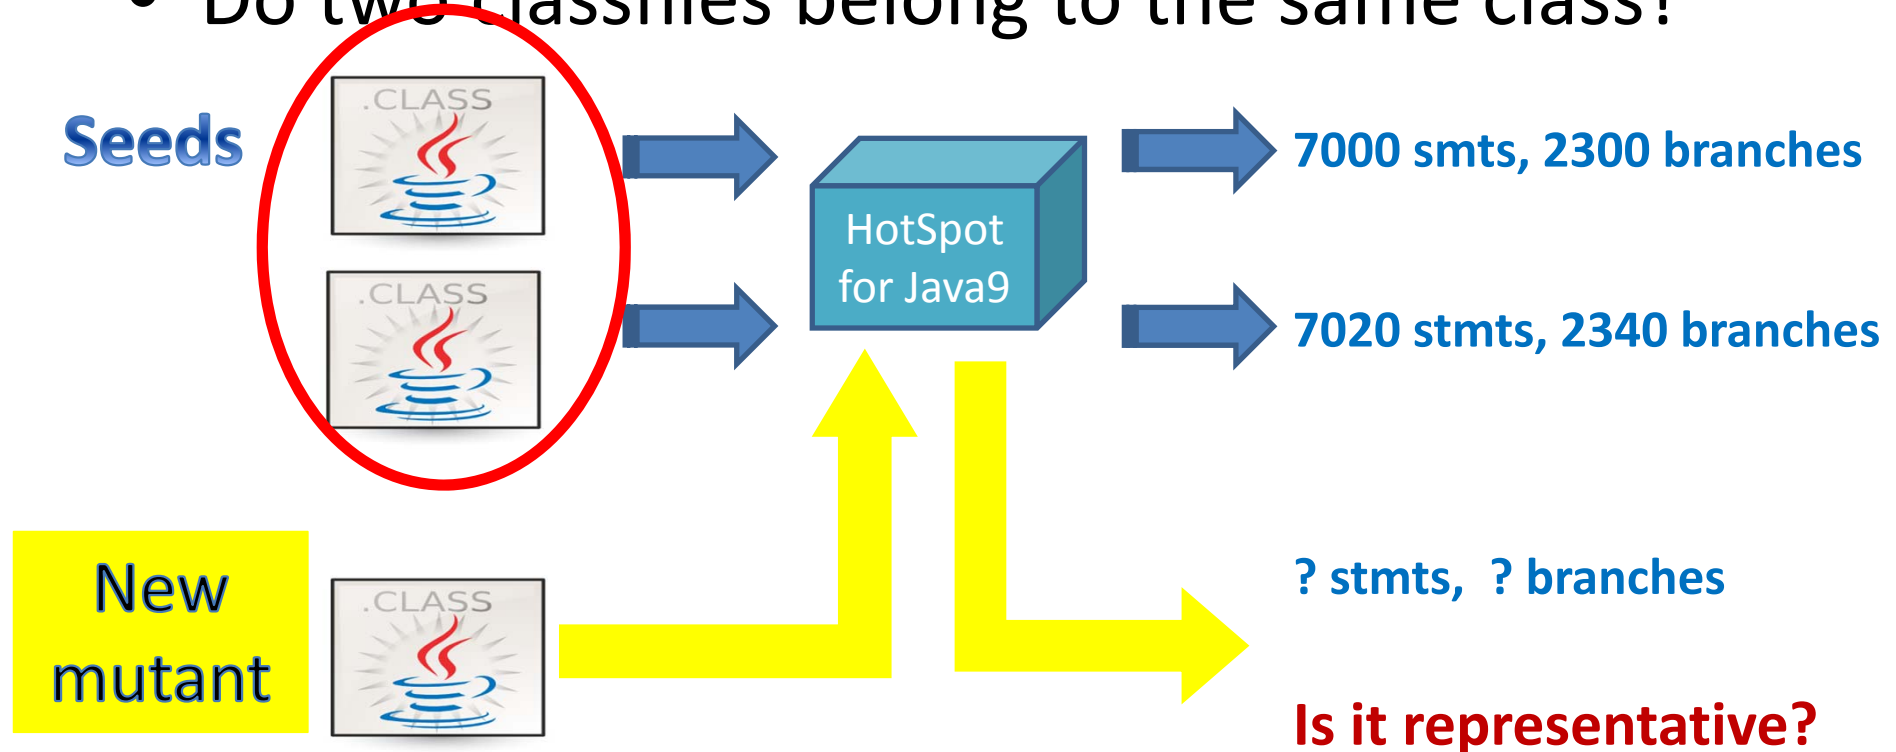
- Do two classfiles belong to the same class?





Selecting Representative Classfiles (2)

- Do two classfiles belong to the same class?



Several comparison criteria can be given here


Outline

- Motivation
- Design
 - Mutating classfiles
 - Selecting representative classfile mutants
 - Selectively applying mutators
 - Differentially testing JVMs
- Evaluation
- Related Work and Conclusion

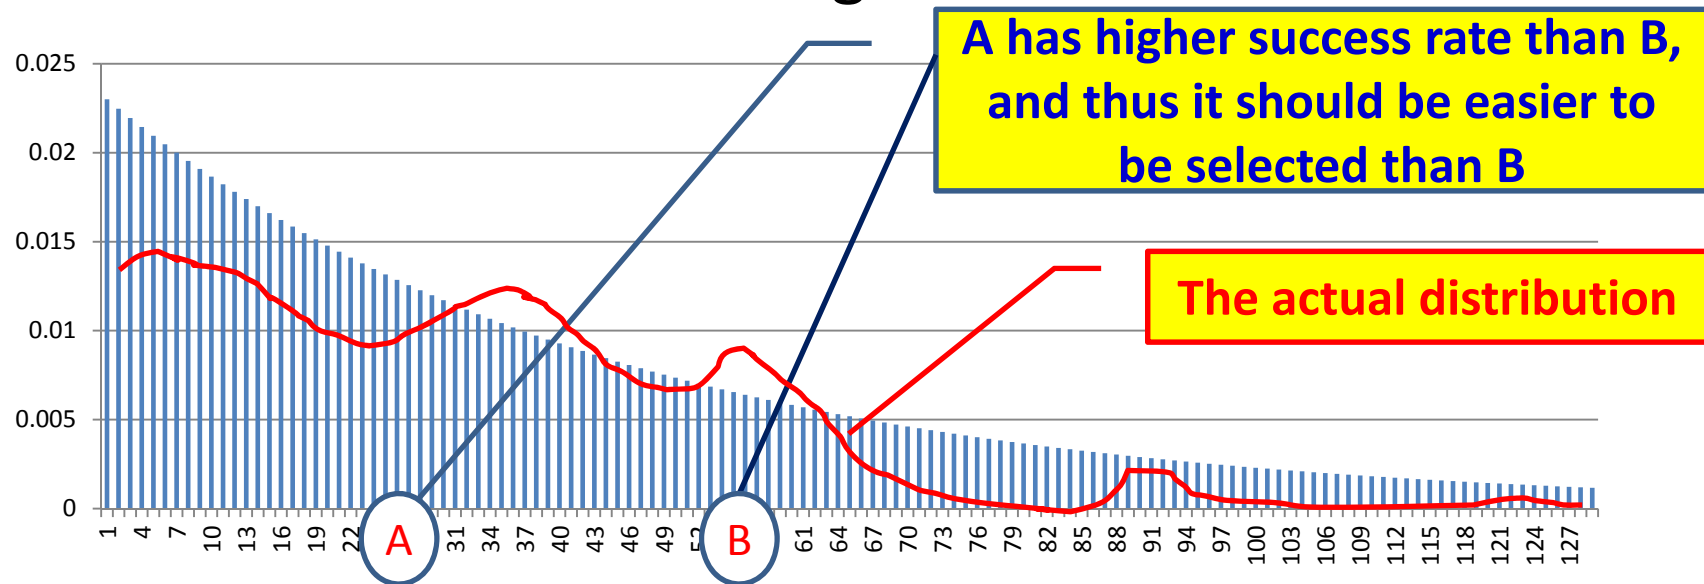
Selecting Mutators

- **Goal:** to create as many representative classfiles as possible
- **Fact:** mutators are designed arbitrarily; some are effective, while some others are useless
- **A naïve solution:** to select mutators by learning from prior knowledge

An MCMC Sampling Method



- Which mutator will be selected at each step?
 - A desired distribution: geometric distribution



Proposition: The more number of representative classfiles have been created by a mutator, the more likely the mutator should be selected for further mutations

More Details

- The desired distribution

$$\Pr(X = k) = (1 - p)^{k-1} p$$


- Classfuzz picks up mutators at random, and then accepts or rejects the mutators by a Metropolis choice

$$\begin{aligned} A(mu_1 \rightarrow mu_2) &= \min\left(1, \frac{\Pr(mu_2)}{\Pr(mu_1)}\right) \\ &= \min\left(1, (1 - p)^{k_2 - k_1}\right) \end{aligned}$$

Outline

- Motivation
- Design
 - Mutating classfiles
 - Selecting representative classfile mutants
 - Selectively applying mutators
 - Differentially testing JVMs
- Evaluation
- Related Work and Conclusion

Execution Comparison



$\text{result}_0 = \text{jvm}_0(\text{env}_0, c, \text{input})$

$\text{result}_1 = \text{jvm}_1(\text{env}_1, c, \text{input})$


$\text{result}_2 = \text{jvm}_2(\text{env}_2, c, \text{input})$

$\text{result}_3 = \text{jvm}_3(\text{env}_3, c, \text{input})$

$\text{result}_4 = \text{jvm}_4(\text{env}_4, c, \text{input})$

A *JVM discrepancy* appears when $\text{result}_i \neq \text{result}_j$
It can either be a **JVM defect** or a **compatibility issue**

Execution Comparison (2)



$\text{result}_0 = \text{jvm}_0 (\text{env}_0, c, \text{input})$

$\text{result}_1 = \text{jvm}_1 (\text{env}_0, c, \text{input})$


$\text{result}_2 = \text{jvm}_2 (\text{env}_0, c, \text{input})$

$\text{result}_3 = \text{jvm}_3 (\text{env}_0, c, \text{input})$

$\text{result}_4 = \text{jvm}_4 (\text{env}_0, c, \text{input})$

A *JVM defect* appears when $\text{result}_i \neq \text{result}_j$

Execution Comparison (3)



$\text{result}_0 = \text{jvm}_0 (\text{env}_0, c, \text{input})$

$\text{result}_1 = \text{jvm}_1 (\text{env}_0, c, \text{input})$

$\text{result}_2 = \text{jvm}_2 (\text{env}_0, c, \text{input})$

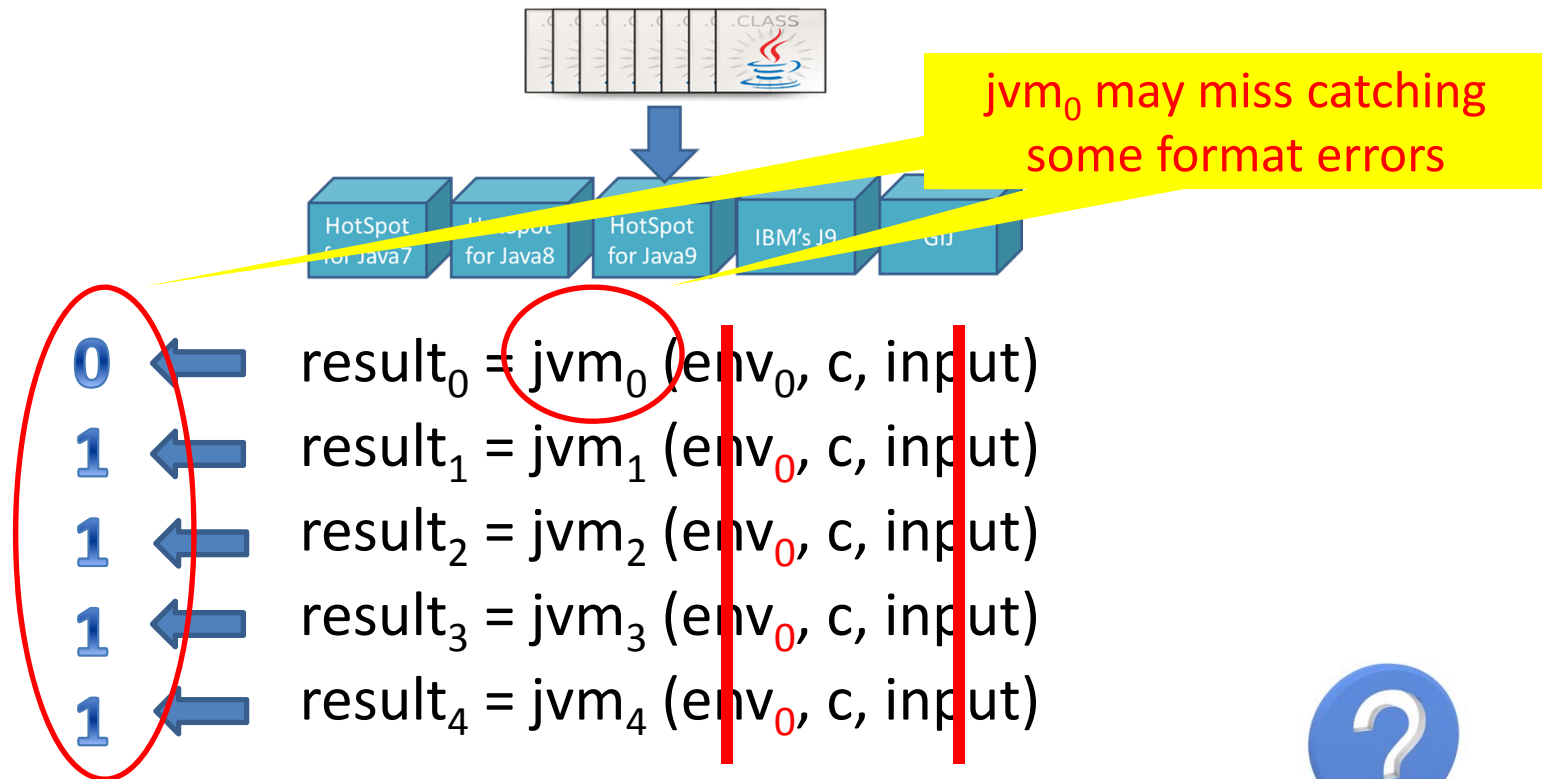
$\text{result}_3 = \text{jvm}_3 (\text{env}_0, c, \text{input})$

$\text{result}_4 = \text{jvm}_4 (\text{env}_0, c, \text{input})$

A *JVM defect* appears when $\text{result}_i \neq \text{result}_j$



Execution Comparison (3)



A *JVM defect* appears when result_i ≠ result_j

Outline

- Motivation
- Design
- Evaluation
 - Setup
 - Results
- Related Work and Conclusion

Setup

- Coverage collection

- HotSpot for Java9
- GCOV + LCOV



At each run the coverage can be conveniently collected

HotSpot (260K LOCs)
Cost for cov. analysis: 30+ mins

share/vm/classfile/ (11977 LOCs)
Cost for cov. analysis: 90 secs

- Seeds

- 1216 classfiles in JRE 7

Evaluated Methods

- Classfuzz supplemented with a uniqueness criterion
 - [st], [stbr], [tr] – explained in the paper
- Randfuzz, Greedyfuzz, Uniquefuzz

	classfuzz			randfuzz	greedyfuzz	uniquefuzz
Mutation-based	✓			✓	✓	✓
Cov. analysis	✓			✗	✓	✓
Uniqueness criterion	[st]	[stbr]	[tr]	✗	[stbr]	[stbr]
Mutator selection	✓			✗	✗	✗


Metrics

- **RQ1:** How **many** test classfiles can be generated?
 - *#Iterations, |GenClasses|, |TestClasses|*
- **RQ2:** How **effective** are the test classfiles?
 - *|Discrepancies|, |Distinct Discrepancies|, diff rate*
- **RQ3:** Can the test classfiles find any JVM **defects**?


Outline

- Motivation
- Design
- Evaluation
 - Setup
 - Results
- Related Work and Conclusion

Results on Classfile Generation



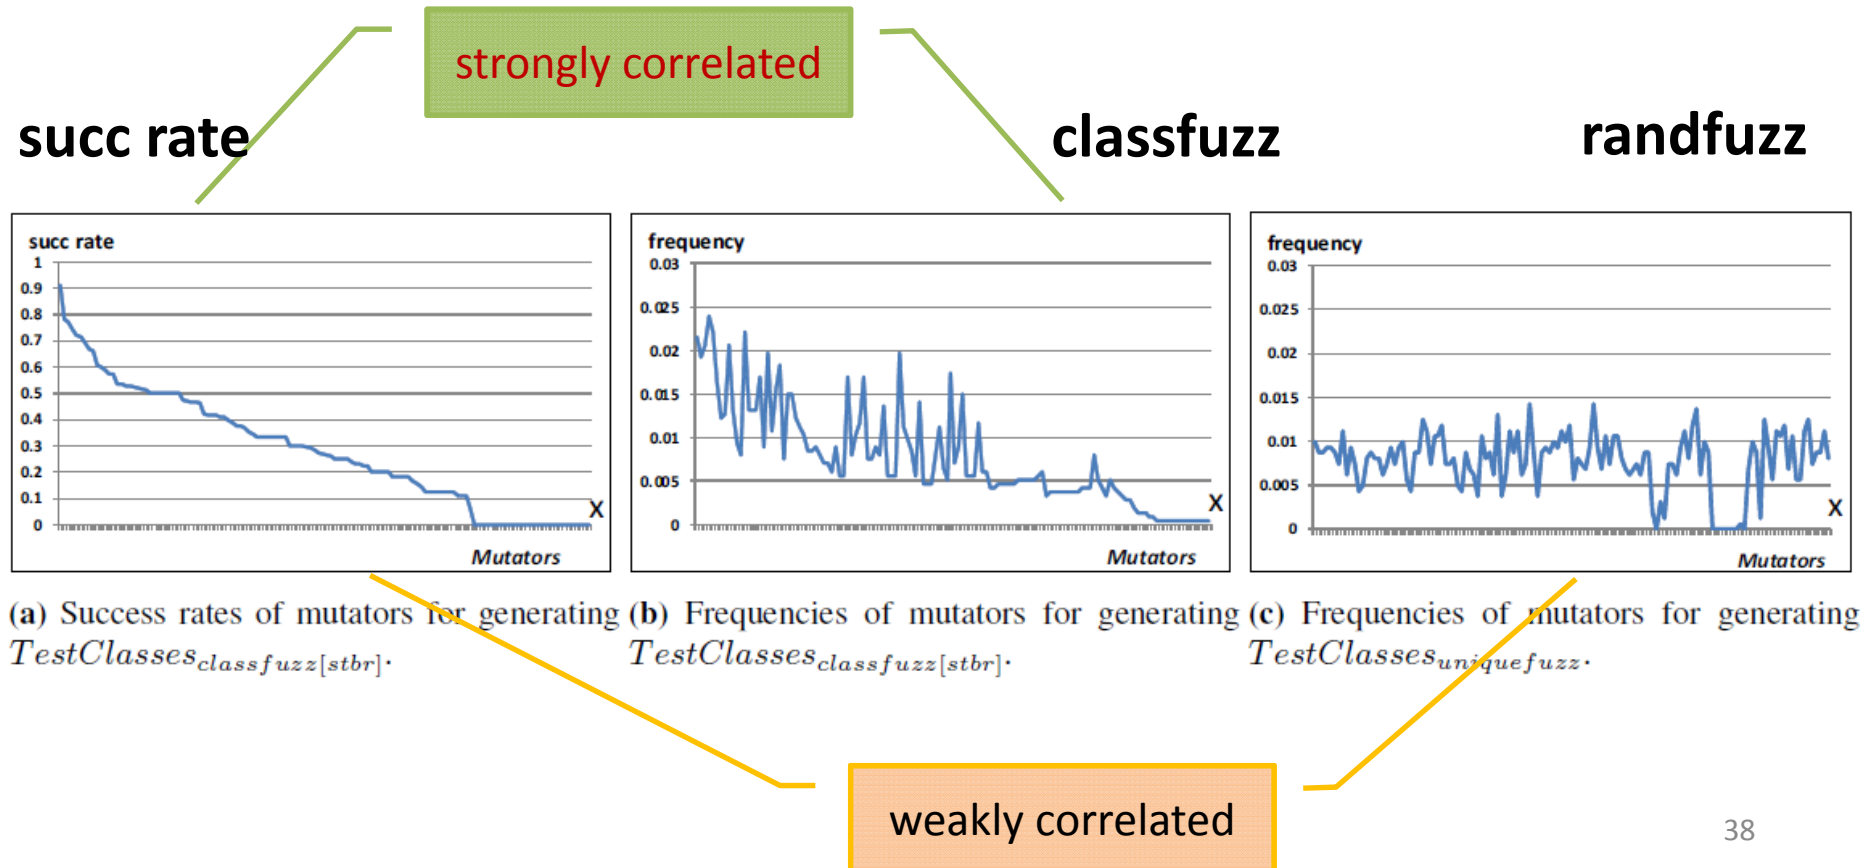
- **Randfuzz generates 20 times as many classfiles as those generated by any other algorithm**
- **Classfuzz[stbr] generates the most number of representative classfiles**



Classfuzz[stbr] achieves the highest success rate among all the coverage-directed algorithms


Results on Classfile Generation (2)

- Classfuzz can utilize the prior knowledge to select mutators



Results on Differential JVM Testing

- Classfuzz can enhance the ratio of discrepancy triggering classfiles from 1.7% to 11.9%



- JVMs are compatible for **most** of the classfiles, but differ in processing corner cases
- We have experienced **898** different execution paths. **107** paths were related to JVM behavior differences

Discrepancy Analysis (1)

```
1  ...
2  MD5 checksum 8fb69050bbcb9a83ddd90ae393368c5e
3  ...
4  class M1436188543
5  minor version: 0
6  major version: 51
7  flags: ACC_SUPER
8  Constant pool:
9  ...
10 #7 = Utf8 <clinit>
11 #8 = Utf8 ()V
12 #9 = Class #19 // java/lang/System
13 #10 = Utf8 Code
14 #11 = Utf8 main
15 ...
16 {
17   public abstract {};
18   flags: ACC_PUBLIC, ACC_ABSTRACT
19   public static void main(java.lang.String[]);
20   flags: ACC_PUBLIC, ACC_STATIC
21   Code:
22     stack=2, locals=1, args_size=1
23     0: getstatic #12 // Field java/lang/System.out:
        Ljava/io/PrintStream;
24     3: ldc #4 // String Completed!
25     5: invokevirtual #21 // Method java/io/
        PrintStream.println:(Ljava/lang/String;)V
26     8: return
27 }
```

public abstract {};

- HotSpot takes it as a **ordinary** method
- J9 reports a **format error**



The JVM specification needs to be clarified

Discrepancy Analysis (2)

```
//The Jimple code of M1433982529
public class M1433982529 extends java.lang.Object
{
    protected void internalTransform(java.lang.
        String)
    {
        java.util.Map r0;
        r0 := @parameter0: java.util.Map;
        staticinvoke <java.lang.Object: boolean
            getBoolean(java.util.Map)>(r0);
        return;
    }
}
```

A type casting needs to be performed



JVMs take their own classfile verification and type checking policies

Discrepancy Analysis (3)

```
//The source code of sun.java2d.pisces.  
    PiscesRenderingEngine  
public class PiscesRenderingEngine extends  
    RenderingEngine {  
    ...  
    private static enum NormMode {OFF, ON_NO_AA,  
        ON_WITH_AA};  
    ...  
}  
//The Jimple code of M1437121261  
public class M1437121261 {  
    public static void main (String[] r0)  
        throws sun.java2d.pisces.  
            PiscesRenderingEngine$2{  
    ...  
}  
}
```



JVMs are not compatible to access some classes

Discrepancy Analysis (4)

- More findings
 - J9 is less strict than HotSpot because J9 only verifies a method when it is invoked, while HotSpot verifies all methods before execution
 - GJ can execute an interface having a main method
 - GJ accepts a class with duplicate fields
 - ...
- These discrepancies can be found in a package of 62 discrepancy-triggering classes

Outline

- Motivation
- Design
- Evaluation
- Conclusion

Conclusion

- **Problem**
 - Testing JVMs requires painstaking effort in designing **test classfiles** along with their **test oracles**
- **Proposal: classfuzz: coverage-directed fuzz testing**
 - **Test classfile generation**
 - Mutating classes and selectively applying mutators
 - Deciding the representativeness of a classfile mutant
 - **Differential JVM testing**
- **Tool is available at**
<http://stap.sjtu.edu.cn/~chenyt/DTJVM/index.htm>