




Latte: A Language, Compiler, and Runtime for Elegant and Efficient Deep Neural Networks

Leonard Truong^{φ, τ}, Rajkishore Barik^τ, Ehsan Toton^τ, Hai Liu^τ,
Chick Markley^φ, Armando Fox^φ, Tatiana Shpeisman^τ


^τIntel Labs, ^φUC Berkeley

Deep Learning Research is Thriving

Academia



Thanks Ben Zorn




<http://jponttuset.github.io/xkcd-deep-learning/>

Deep Learning 101

Neural networks: family of biologically inspired, machine-learning models


Neurons are organized into **layers**

DNN Architecture/Model: a specific configuration of layers



Layer 1 Layer 2 Layer 3

Weighted Neuron: fundamental building block of neural networks



$$f(x) = b + \sum_i w_i * x_i$$

Parameters $\{w_1, w_2, w_3, b\}$ are learned


Deep Learning 101

Training neural networks

For each training sample

- **Forward propagation**
- **Back propagation**
- Update parameters


Typically process the entire training data set multiple times



Deep Learning Research

New Layers
New Network Architectures

What's the best programming environment for deep learning research?



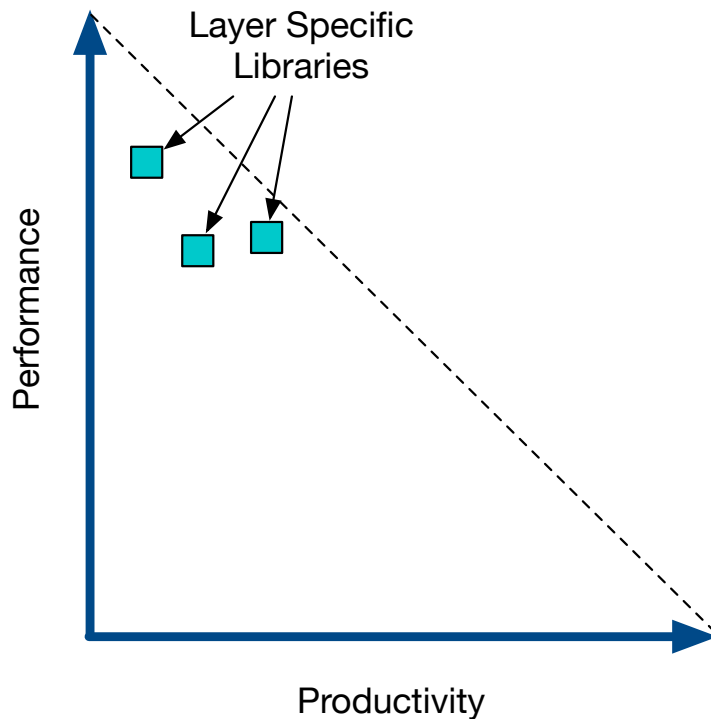
Can take **days to weeks** to evaluate a new idea

- AlexNet 5-6 days [Krizhevsky et al. NIPS '12]
- VGG 2-3 weeks [Simonyan et al. ICLR '15]

Deep Learning Software Landscape

Layer Specific Libraries


- Provide a function for each layer
- **Examples:** cuDNN, Caffe, NNPACK, ...
- Good performance
- Hard to program
- **Cannot fuse across layers**



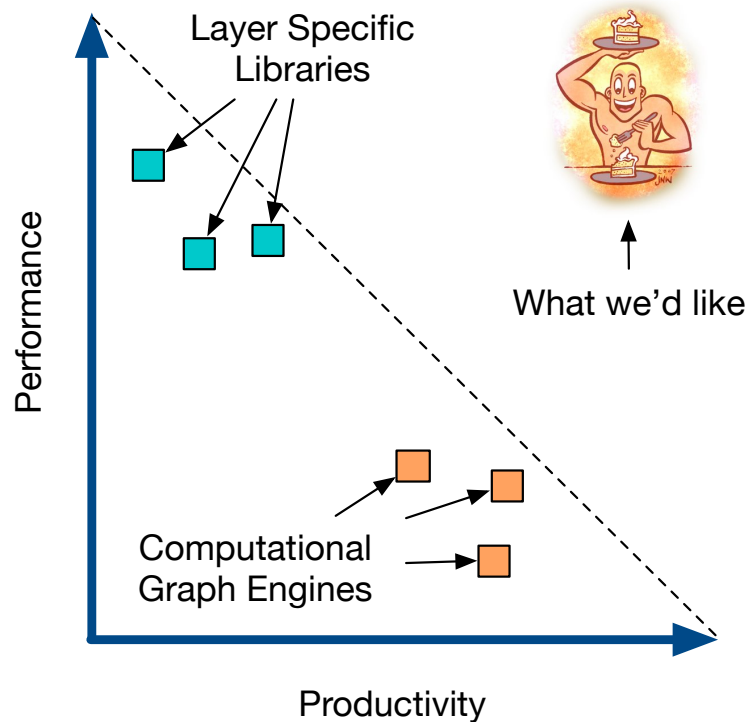
Deep Learning Software Landscape

Computational Graph Engines

- DNNs can be expressed as graphs of general operations
- **Examples:** Theano, CNTK, TensorFlow, ...
- Easier to program
- Sacrifice performance
- **Many rely on bindings to cuDNN**



Deep Learning Software Landscape



Introducing Latte



DSL

- Language for describing DNNs as a system of interconnected neurons

Compiler

- Constructs an implicit data-flow graph from user description
- Synthesizes and optimizes an implementation of the data-flow graph

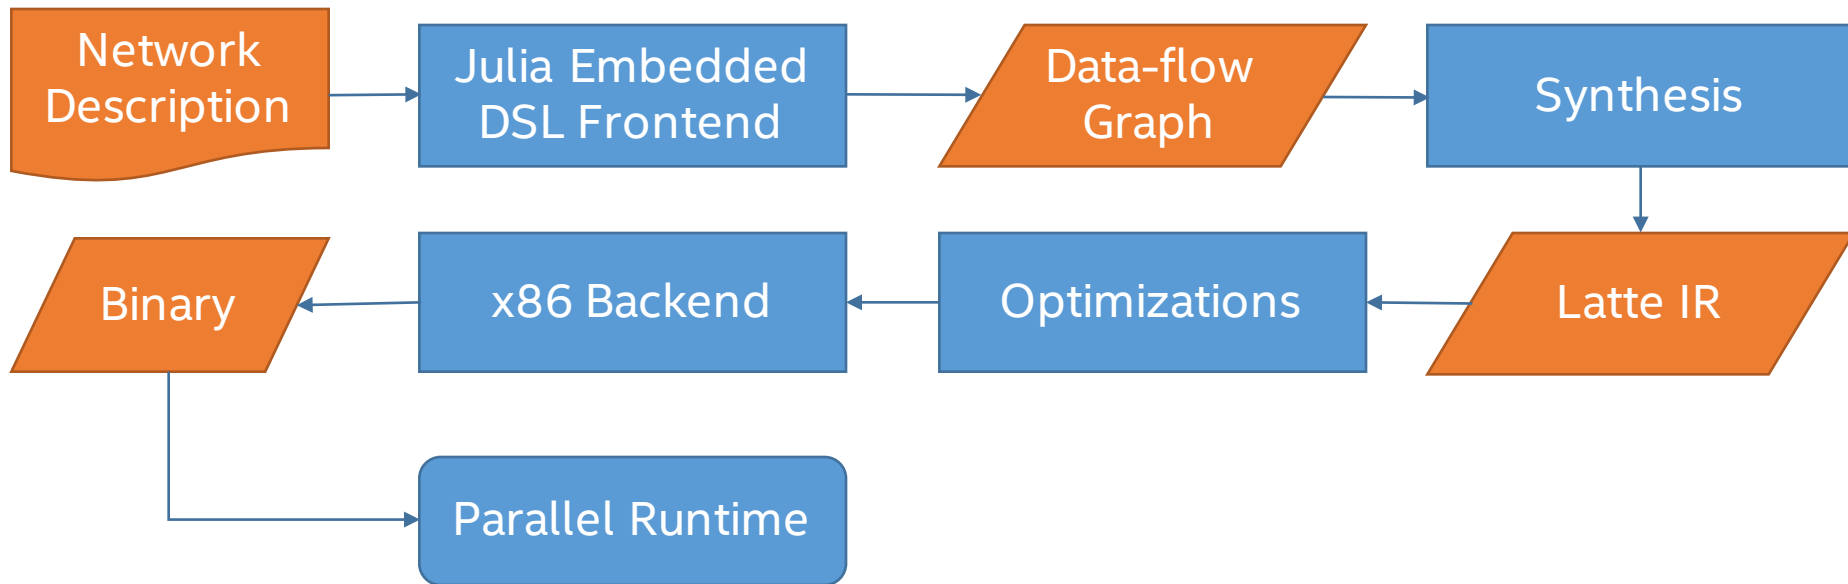
Runtime

- Supports distributed memory parallelism for data-parallel training

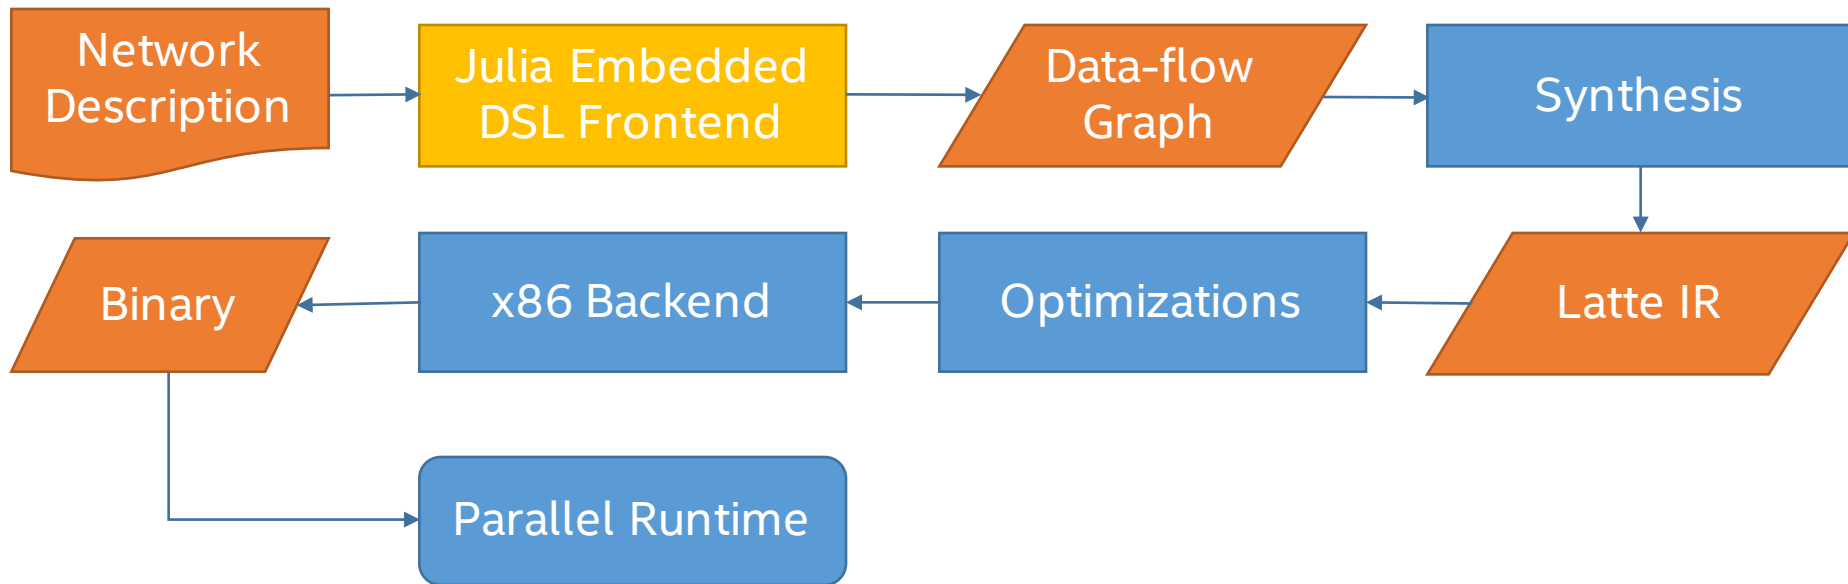
Competitive Performance

- 3-6x speedup over Caffe (C++/MKL) on a single node


Latte System Diagram



Latte DSL Frontend



Latte DSL: Introduction



Language Constructs

 Ensemble

 Neuron

 Connection

Embedded in



Latte DSL: Neuron

Neuron – Primitive, abstract data-type

```
type Neuron
  value      :: Float32
   $\nabla$        :: Float32
  inputs     :: Vector{Float32}
   $\nabla$ inputs  :: Vector{Float32}
end
```


User can

- Define additional fields for internal state (optional)
- Implement ***forward*** and ***backward*** functions (required)

Latte DSL: Neuron Example

WeightedNeuron adds additional fields for learning weights and bias values

```
@neuron type WeightedNeuron <: Neuron
  weights    :: Vector{Float32}
  ∇weights   :: Vector{Float32}
  bias       :: Float32
  ∇bias      :: Float32
end
```




- weights holds $\{w_1, w_2, w_3\}$, ∇ weights holds gradients
- bias holds $\{b\}$, ∇ bias holds gradient
- **Important:** Inherits value, ∇ , inputs, and ∇ inputs from Neuron

Latte DSL: Neuron Forward Example

The forward propagation of a **WeightedNeuron**

$$f(x) = b + \sum_i w_i * x_i$$



```
@neuron forward(neuron::WeightedNeuron) do
  for i in 1:length(neuron.inputs)
    neuron.value += neuron.weights[i] * neuron.inputs[i]
  end
  neuron.value += neuron.bias
end
```

Latte DSL: Ensemble

An **Ensemble** is an **N**-dimensional array of a Neuron subtype **T**

```
type Ensemble{T <: Neuron, N}  
  name      :: Symbol  
  neurons   :: Array{T,N}  
  connections :: Vector{Connection}  
  ...  
end
```




Latte DSL: Connections

Connections between a *source* ensemble and a *sink* ensemble are defined using a **mapping function**

```
add_connections(net, source, sink, (i) -> (i:i+1,))
```

1

1:4




Bringing it all together: Fully Connected Layer

```
# Create a 1-D array of `num_outputs` WeightedNeurons
neurons = [WeightedNeuron(weights[:, i],  $\nabla$ weights[:, i],
                           bias[:, i]    ,  $\nabla$ bias[:, i])
            for i in 1:num_outputs]

# Construct the ensemble
ens = Ensemble(net, name, neurons)

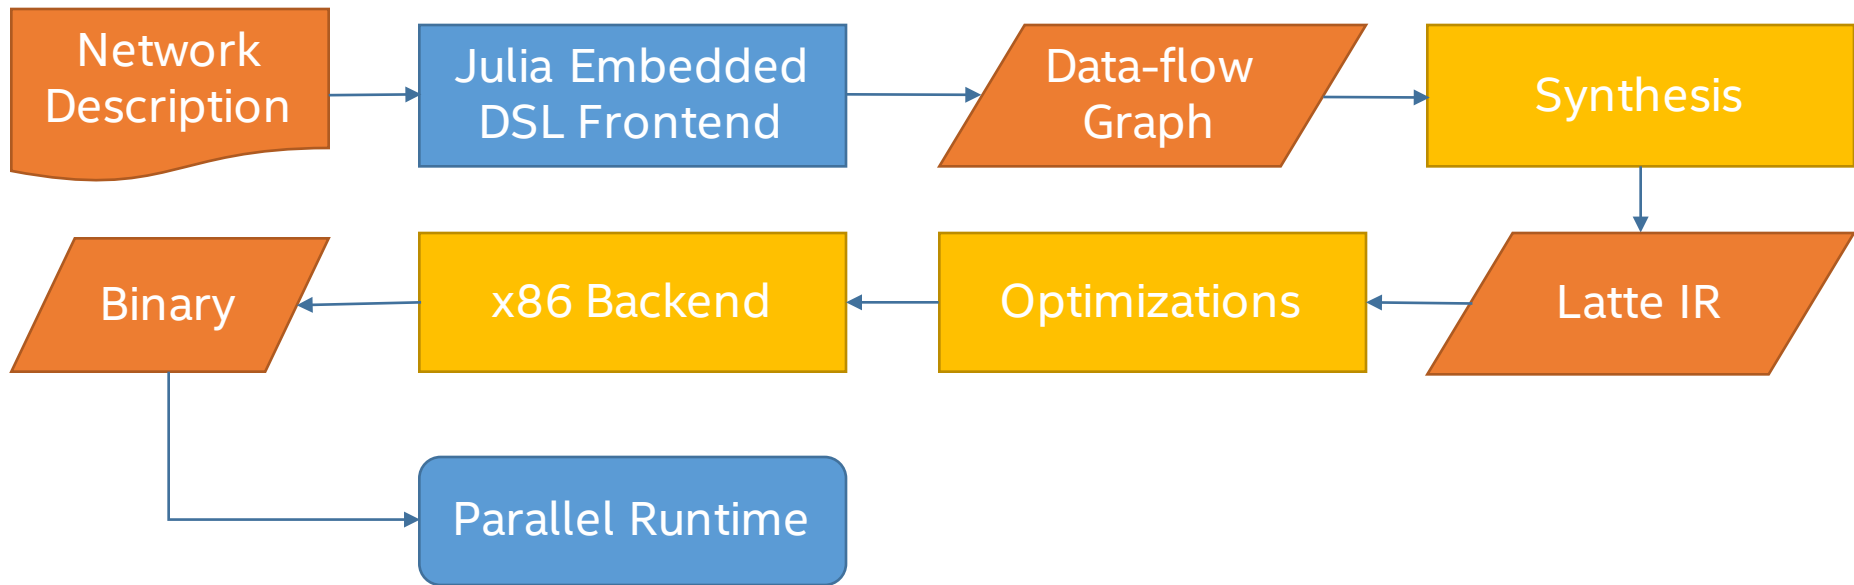
# Connect each neuron in input_ensemble
add_connections(net, input_ensemble, ens,
               (i) -> (1:d for d in size(input_ensemble)))
```



input_ensemble

ens

Latte Compiler



Latte Compiler: Synthesis

User implicitly describes a **data-flow graph (DFG)**

- Stored implicitly by mapping functions
- Ensembles provide a natural partitioning of the DFG

Data-flow Synthesis


- Copy output values of connected neurons into an input buffer for an ensemble

Compute Synthesis

- Perform the computation of neurons in an ensemble

Distributed Memory Communication

- Asynchronously synchronize gradients



Latte Compiler: Optimizations

Cross Layer Fusion (Covered in this talk)

- Use dependence information from data-flow graph to fuse loops across layers

Shared Variable Optimization (See paper)

- Compiler analysis discovers values that are shared between neurons

Library Kernel Pattern Matching (See paper)

- Replace matrix-multiplication loop nest patterns with MKL call

General Loop Optimizations (See paper)

- Tiling, Fusion

Latte Compiler: Cross Layer Fusion Example

Conv
Layer

```
for y_tile in 1:TILE_SIZE:height
    gemm('T', 'N', TILE_SIZE*width, n_filters, n_inputs,
        conv1input[n], conv1weights, conv1[n])
```

Easily Fused

ReLU
Layer

```
for y_tile in 1:TILE_SIZE:height
    for c = 1:n_filters,
        y = y_tile:y_tile+TILE_SIZE,
        x = 1:width
        conv1[x, y, c] = max(conv1[x, y, c], 0.0)
```

Different Loop Lengths

Pooling
Layer

```
for y_tile in 1:TILE_SIZE:height/2
    for c = 1:n_filters,
        y = y = y_tile:y_tile+TILE_SIZE,
        x = 1:width
        for p = 1:2, q = 1:2
            poolinput[p*2+q, x, y, c] = conv1[x+q, y+p, c]
            maxval = -Inf
            for i = 1:pool_size
                maxval = max(poolinput[i, x, y, c], maxval)
            pool1[x, y, c] = maxval
```

Possible Fusion Preventing
Dependency

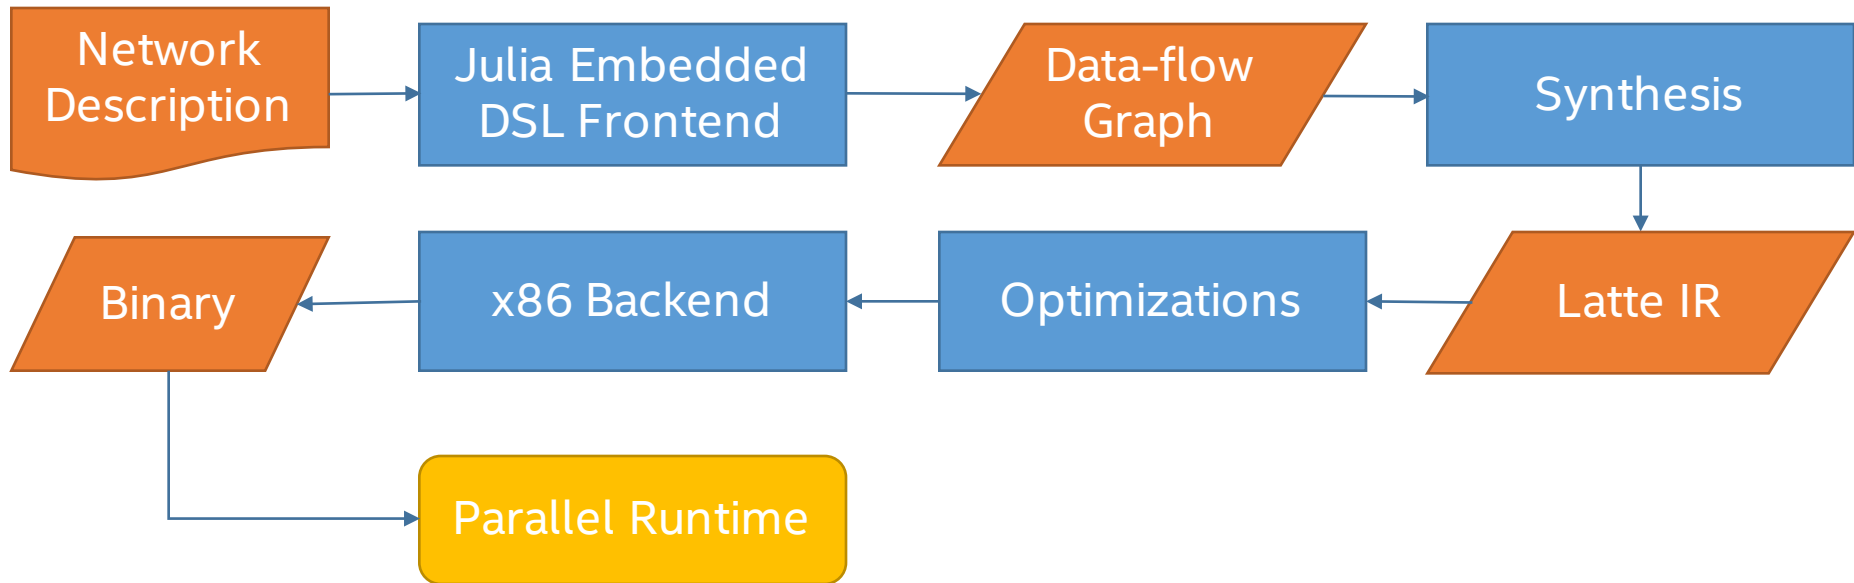
Latte Compiler: Cross Layer Fusion Example

Key Takeaways

- Fusing across layers is only possible in a dynamic compiler, **static libraries cannot do this**
- Dependence information already provided by data-flow graph

```
for y_tile in 1:TILE_SIZE:height/2
    gemm('T', 'N', TILE_SIZE*2*width,
        n_filters, n_inputs, conv1input[n],
        conv1weights, conv1[n])
    for c = 1:n_filters,
        for y = y_tile:y_tile+TILE_SIZE*2,
            x = 1:width
                conv1[x, y, c] = max(conv1[x, y, c],
                                     0.0)
        for y = y_tile:y_tile+TILE_SIZE,
            x = 1:width
                maxval = -Inf
                for p = 1:2, q = 1:2
                    maxval = max(conv1[x+q, y+p, c],
                                maxval)
                pool1[x, y, c] = maxval
```

Latte Runtime




Latte Runtime

Manages data-parallel training on distributed memory systems

- Gradients reduction is overlapped with forward/backward phases

Supports automatic offloading to accelerators (Xeon Phi)

- Data movement to and from accelerator is overlapped with host compute



Evaluating Latte

Compared performance against open-source frameworks

- Caffe [C++/MKL], Mocha.jl [Julia/MKL] (see paper for Mocha results)
- **Training throughput** used as metric

Platforms in this talk:


- Single node : Dual-socket Intel Xeon E5-2699 v3 (36 cores)
- Supercomputer : NERSC Cori Phase 1 (info at nersc.gov)

More Platforms (see paper):

- Accelerator Offloading (Xeon + Xeon Phi), Commodity Cluster

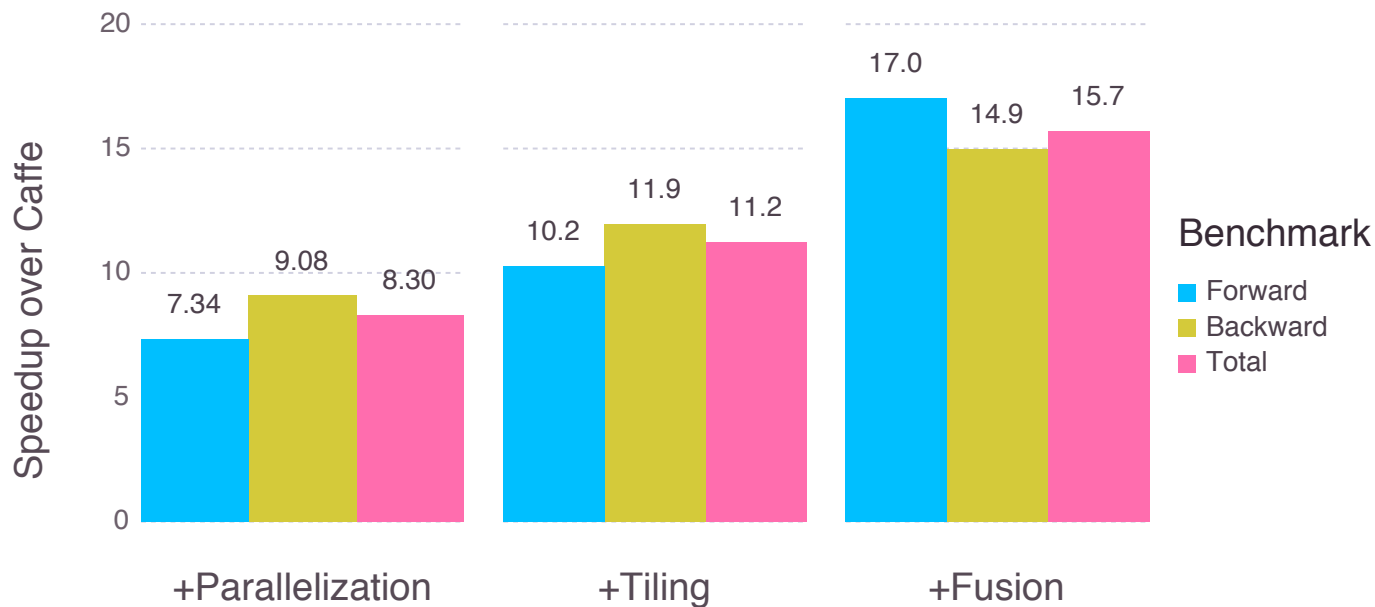
Latte Single Node Evaluation

3-6x speedup over Caffe for Training (Forward + Backward)



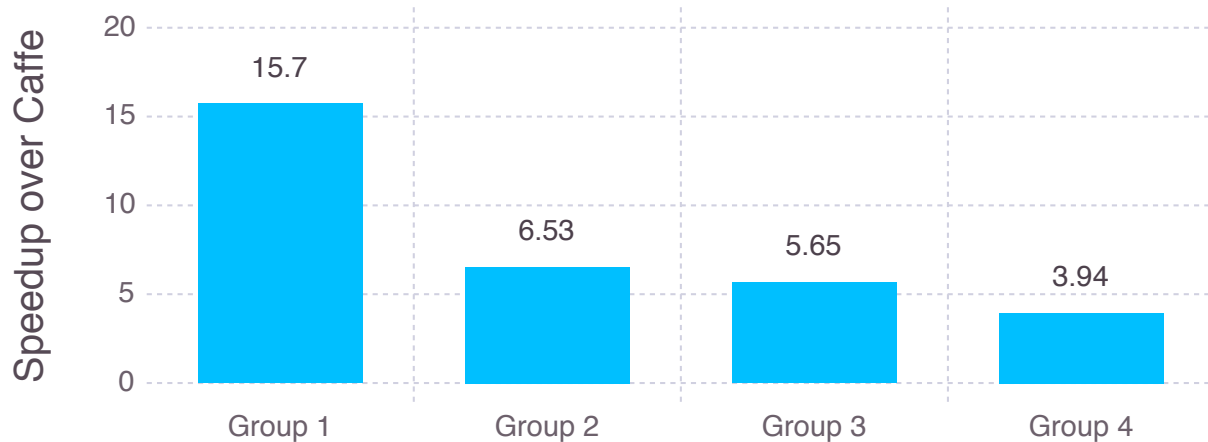
Latte Single Node Evaluation

Breakdown of optimization effectiveness for first 3 layers of VGG



Latte Single Node Evaluation

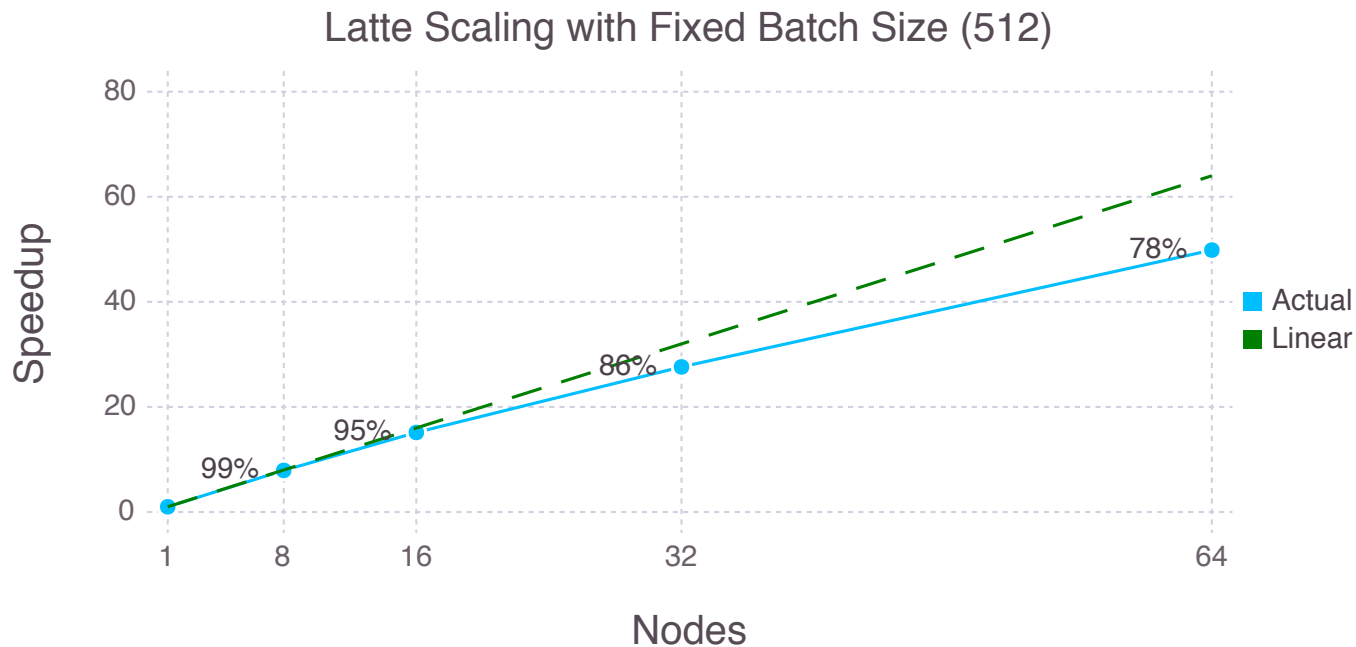
Breakdown of speedup for VGG (Forward)



First 16 layers of VGG

VGG	
Input (128, 3, 224, 224)	Group 1
Conv filt=64, kernel=3, pad=1	
ReLU	
Pool kernel=2, stride=2	
Conv filt=128, kernel=3, pad=1	Group 2
ReLU	
Pool kernel=2, stride=2	Group 3
Conv filt=256, kernel=3, pad=1	
ReLU	
Conv filt=256, kernel=3, pad=1	
ReLU	Group 4
Pool kernel=2, stride=2	
Conv filt=512, kernel=3, pad=1	
ReLU	
Conv filt=512, kernel=3, pad=1	
ReLU	
Pool kernel=2, stride=2	

Latte on Cori



Each worker is given a $512 / N$ sized sub-batch (N = number of Nodes)

Conclusion

Developed a novel abstraction for describing neural networks

Demonstrated performance competitive with existing frameworks

Code: Latte.jl is available at <https://github.com/IntelLabs/Latte.jl>

Ongoing Work: Forthcoming release of Python implementation with new features and better performance



Acknowledgements



Thanks to **Brian Lewis**, **Lindsey Kuper**, and **Justin Gottschlich** for their valuable feedback on the paper and presentations.

Thanks to **Paul Hargrove** at NERSC for access to Cori

For the UC Berkeley contributors, research was partially funded by DARPA Award Number HR0011- 12-2-0016, the Center for Future Architecture Research, a member of STARnet, a Semiconductor Research Corporation program sponsored by MARCO and DARPA, and AS- PIRE Lab industrial sponsors and affiliates Intel, Google, Hewlett-Packard, Huawei, LGE, NVIDIA, Oracle, and Samsung.