

June 17, 2016 - PLDI

# Assessing the Limits of Program-Specific GC Performance

Nicholas Jacek, Meng-Chieh Chiu,  
Ben Marlin, and Eliot Moss

{njacek,joechiu,marlin,moss}@cs.umass.edu  
University of Massachusetts, Amherst

# What this talk is about

- Garbage Collection (GC)
- Lower bounds on GC cost
- Not generic, but for a particular program run
- After-the-fact analysis – not a mechanism
- Optimal for full-heap GC
- Approximately optimal for generational GC
- Optimization methods from machine learning
- Some tricks to make optimization practical


# Program-Specific GC

- Existing GC performance bounds framed in terms of best a GC algorithm can do in the face of *any possible* program behavior
  - Argued by devising an adversarial program
- If we tune GC to a program, or even to a program run, the problem is different:
  - What is the best we can do in the face of a particular sequence of allocations and object deaths, given heap space  $S$  ?

# Why is this interesting?


- In our larger research project we aim to tune GC for individual programs.
  - How do we know how well we are doing?
  - Suppose we can indicate  $x\%$  improvement over some existing scheme. Is there more to be had?
  - If we fail to see improvement, could it be because very little is possible?
- Determines whether, and maybe for which programs, this tuning might be interesting.

# Model of GC



[enabling predicate] Action(parameters)

# Heap states, action effects



[enabling predicate] Action(parameters): # of bytes {cost}

- Invariant:  $F + L + D = S$  (heap size)
- Actions change the state of bytes
- Visible to controller:  $F$  (*not*  $L$  or  $D$ )

# Cost model and optimization


- Cost of a GC taken as proportional to  $L$  (live bytes)
  - More or less true for tracing / copying collectors
- Since the live size at any given point during a program's execution is a fixed property of the execution, cost to collect at time  $t$  is a constant
  - The Markov Decision Process is *first-order* – does not depend on the history of prior decisions
- Therefore, dynamic programming will find an exact solution to determining the optimal GC schedule – places to collect to obtain minimal total cost – for a particular trace (program execution)

# Solution Cost

- $N$  = number of units (objects) allocated
- Cost to solve is  $O(N^2)$ 
  - Assuming you have the live size at each point
- Can refine to  $O(H \cdot N)$  where  $H$  is the heap size
  - Can look back/forward at most  $O(H)$  allocations



# Cost reduction: Full GC



See the paper for more Full GC results

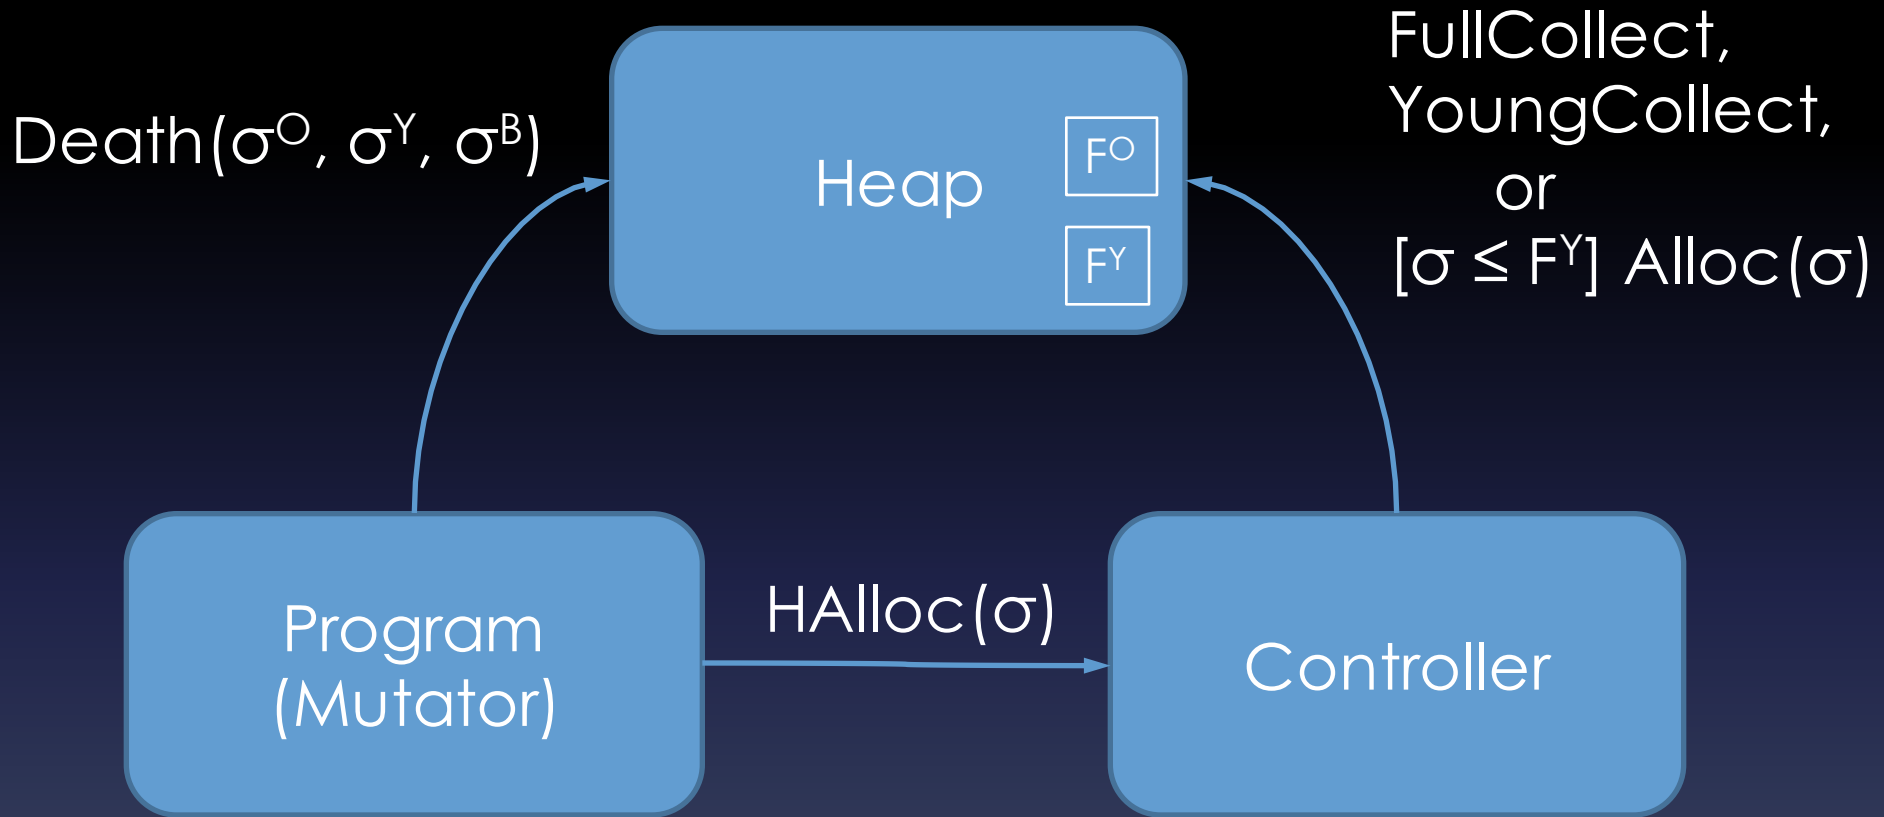
# Making optimization practical

- Group allocations into blocks, say 256 Kb
  - Reduces N by a factor of (say) 5000
  - Does this by constraining when GC occurs
  - Smaller blocks do not change things much
- Pre-analyze object connectivity
  - Identify objects treated the same by GC
  - Summarize behavior in three numbers
  - No detailed simulation while optimizing!

# Generational Collection

- Splits heap into *young* and *old* portions
- Allocation goes into the young generation
- Full collection finds liveness of *all* objects
- Young collection assumes old objects live
  - Thus treats some dead young objects as if live
  - We call these baggage
- Young collection promotes apparently live young objects to old space

# Model with Gen GC



# States, effects with Gen GC

YC:  $L^Y + B$   $\{L^Y + B\}$

FC:  $L^Y$   $\{L^Y\}$

$D(\sigma^O, \sigma^Y, \sigma^B): \sigma^O$

$F^O$  (free old)  
(Initially  $S^O$ )

$L^O$  (live old)  
(Initially 0)

$D^O$  (dead old)  
(Initially 0)

FC:  $D^O$   $\{L^O\}$

$[\sigma \leq F^Y \ \& \ S^Y - F^Y \leq F^O]$

$\text{Alloc}(\sigma): \sigma$

YC, FC: B

B (baggage)  
(Initially 0)

$D(\sigma^O, \sigma^Y, \sigma^B): \sigma^B$

$D(\sigma^O, \sigma^Y, \sigma^B): \sigma^Y$

$F^Y$  (free young)  
(Initially  $S^Y$ )

$L^Y$  (live young)  
(Initially 0)

$D^Y$  (dead young)  
(Initially 0)

YC, FC:  $L^Y$

YC, FC:  $D^Y$


# Impact on Optimization

- Full GC cost, and resulting state, depends only on time of collection
- This is not true for Young GC!
  - Previous promotions affect whether a dying object goes to  $D^Y$  or B ...
  - Which later affects cost (and future promotion)
- We *approximate* by considering visible states based on current time and time of previous collection

# Computing Cost Efficiently

- For a given trace, can precompute necessary reachability information, avoiding simulation
- We have (birth, death) for each object
  - Precise death comes from Elephant Tracks' analysis
- Add prebirth, giving (prebirth, birth, death)
- GC of object in Young space at  $t > \text{death}$  causes promotion if previous GC was between prebirth and birth
- Aggregate blocks of objects by (p,b,d)

# Prebirth Time






# Gen GC Optimization

- Develop collection of sample states:
  - $(F^Y, F^O, t_{\text{last}}, t_{\text{now}})$  where  $t_{\text{last}}$  is time of most recent GC
  - Cost of legal actions (noC, YC, FC) in each state
  - Use “collect when full” to get to each  $t_{\text{last}}$
  - Only a sample – full search space *huge!*
- Treat as Markov decision process
  - Conflation of states treated as randomness
- Apply Least Squares Policy Iteration


# Experiments

- 7 programs from DaCapo suite + javac
- For most of them, added more inputs
- About  $10^2$  to  $10^4$  256 Kb blocks
- Ran under Elephant Tracks to generate traces of alloc, death, pointer updates
- Computed (prebirth,birth,death) times
- Aggregated into blocks
- Applied LSPI
- Compared with default: collect when full

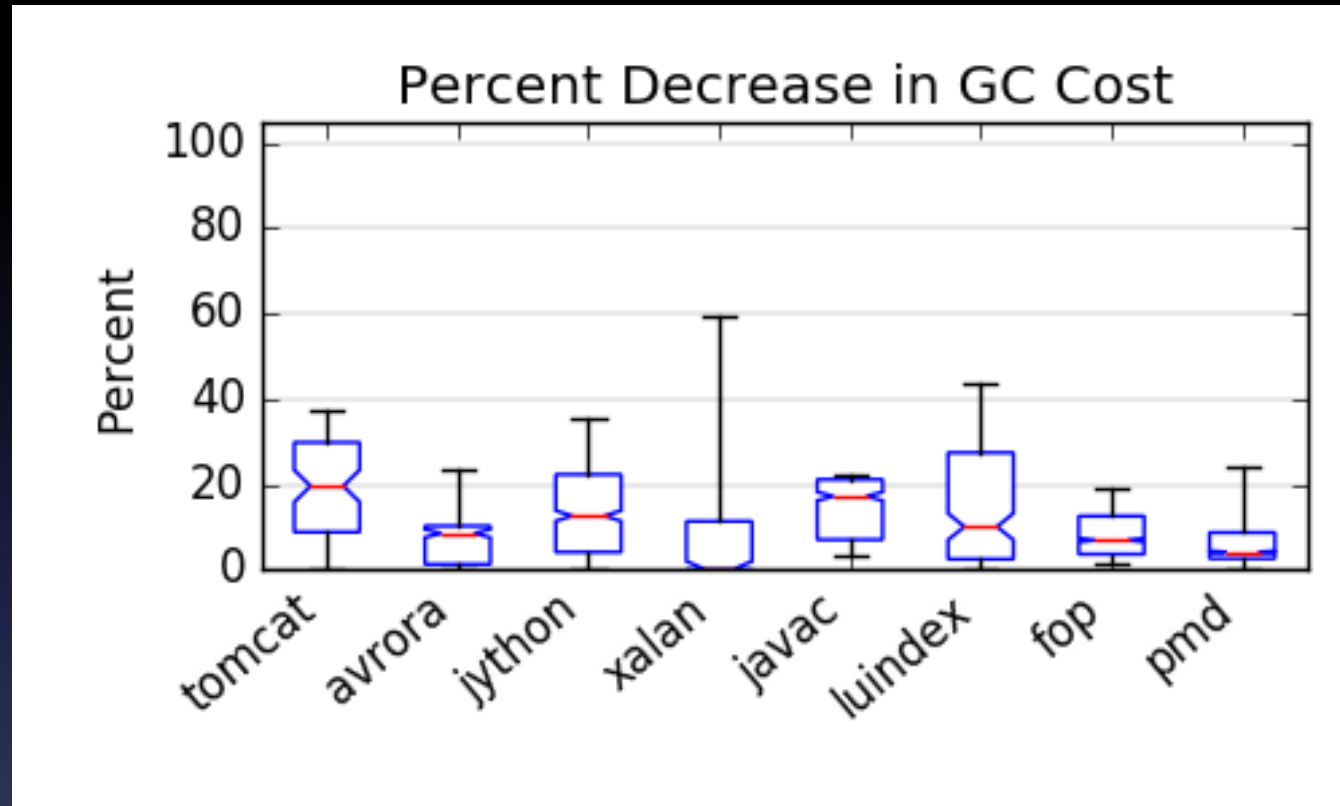
# Results: Sample Schedules



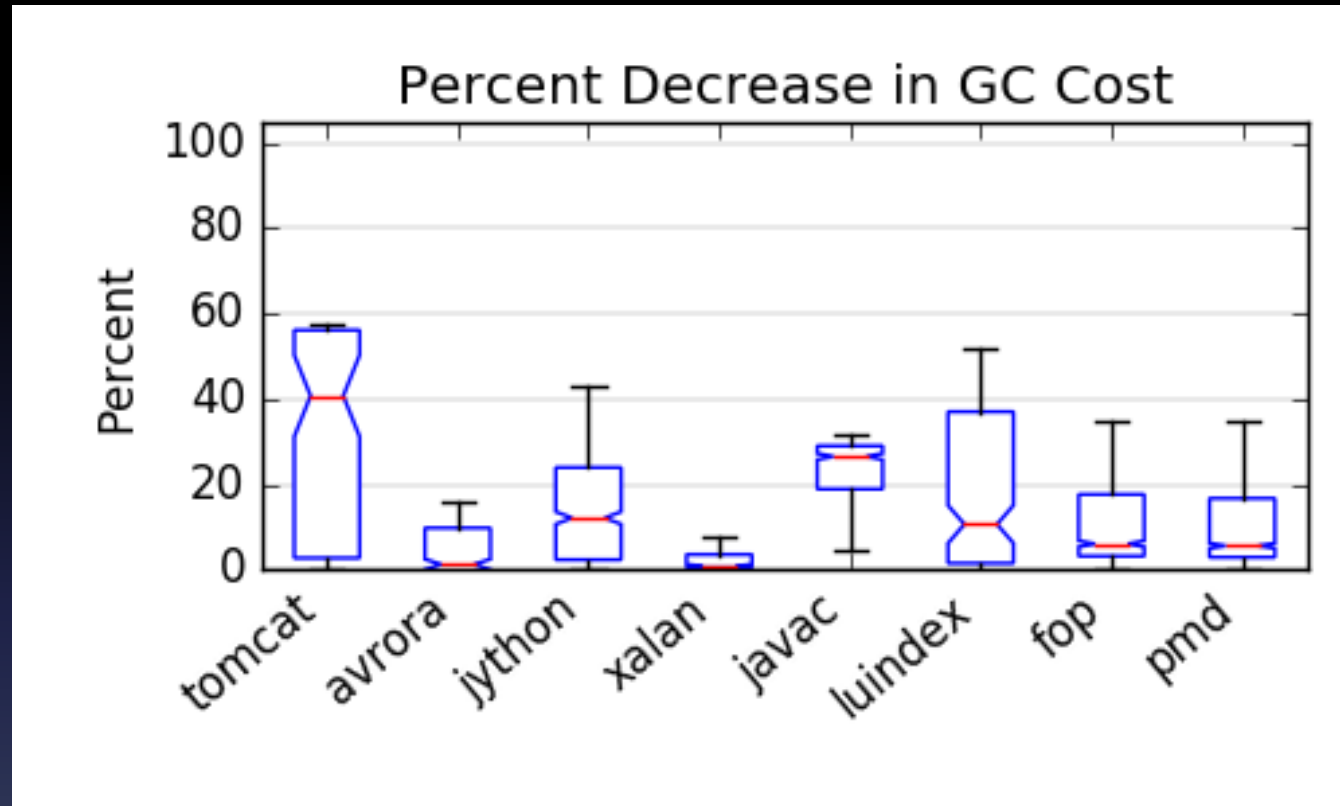
# Cost vs Heap size ( $S^Y=8\text{Mb}$ )



# Reduction: Gen GC ( $S^Y=4\text{Mb}$ )



# Reduction: Gen GC ( $S^Y=8\text{Mb}$ )



# Additional Results

- Improvement is not better for smaller blocks – more fragmentation, etc.
- Cost of optimization is  $\sim O(N^{2.2})$
- More graphs in paper

# Ongoing Work

- Program-specific policies
  - Reinforcement learning, or “deep” learning
    - Collect feature-rich traces
    - Select features
    - Train GC triggering mechanism
  - Some success for “self test”
  - Generalization (over runs, heap sizes) hard
- Hard lower bounds
  - Bound the baggage
  - Solve with dynamic programming



# Conclusions

- First per-program GC cost bounds
- Program-specific GC policies promising
  - For at least some programs and heap sizes
  - Possibly useful reduction in GC time, or
  - Possibly substantial reduction in space
- Sometimes hard to improve over default
- Hard, not approximate, lower bounds would be welcome

# Credits

- Nicholas Jacek – RL work
- Meng-Chieh (Joe) Chiu – ET work
- Ben Marlin – co-PI
  
- NSF Grant CCF-1320498

# Markov Decision Process

- A set of *states*,  $S$
- A set of *actions*,  $A$
- Probabilistic transition function:



# Reinforcement Learning

- States  $s \in S$ , actions  $a \in A$
- Transition function  $T: S \times A \rightarrow S$
- “Policy”  $\pi: S \rightarrow A$  (chooses action)
- Bellman equation: value of each action in each state, for given policy  $\pi$ :

$$Q^\pi(s, a) = c_0 + Q^\pi(s', \pi(s')), \text{ for } s' = T(s, a)$$

In matrix form:

$$Q^\pi = C + \Pi^\pi Q^\pi$$

- Given  $C$  and  $\Pi^\pi$  can solve for  $Q^\pi$

# Policy Iteration (PI)

- Sequence of policies  $\pi_0, \pi_1, \dots$
- Monotonically improving cost
- Determine  $Q^{\pi_m}$  for  $\pi_m$
- Form  $\pi_{m+1}(s) = \arg \min_a Q^{\pi_m}(s, a)$ 
  - That is, lowest cost action in each state, where remaining decision are as for  $\pi_m$
  - This is no worse than  $\pi_m$
- Iterate until reach convergence:  $\pi_*$

# Least Squares PI

- Approximate  $Q^\pi$  as linear combination of a fixed set of basis functions  $\Phi_i(s,a)$ 
  - Weights  $\theta_i$
  - Matrix form:  $Q^\pi = \Phi \theta^\pi$
- Solves with standard linear methods
- We use one basis function for each  $(s,a)$ 
  - So: exact, not (additional) approximation
- Cost is  $O(N^3)$  per iteration, worst case
  - $O(N^{2.2})$  in practice
  - Using blocks was an important choice!

# Learning “Take Home”

- Use a big sample of states and cost
  - Approximation comes in this sampling
  - Our samples not random, but not full space
- Solve with standard linear methods
- Cost is  $O(N^{2.2})$
- Not really learning – an optimization technique borrowed from RL