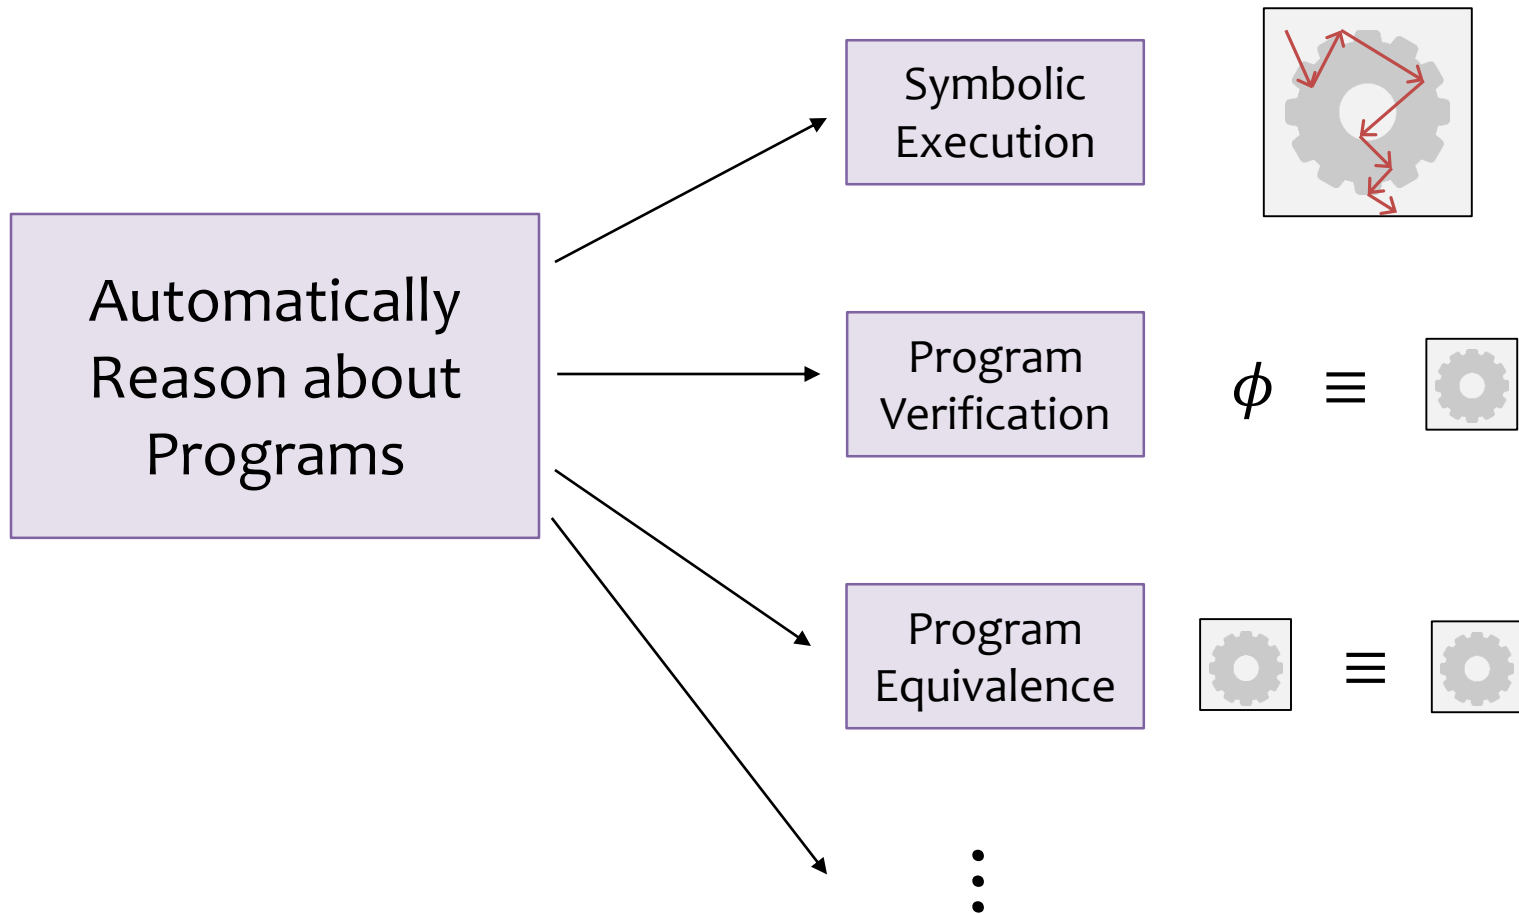


Stratified Synthesis: Automatically Learning the x86-64 Instruction Set

Stefan Heule, Eric Schkufza, Rahul Sharma, Alex Aiken




Motivation



Automatically reasoning about programs requires

Formal Semantics

x86-64



```
testq %rdi, %rdi
je .L1
xorq %rax, %rax
.L0:
movq %rdi, %rdx
andq $0x1, %rdx
addq %rdx, %rax
shrq $0x1, %rdi
jne .L0
cltq
retq
.L1:
xorq %rax, %rax
retq
```

x86-64 Semantics

addq \$0x1, %rax

64-bit bit-vector addition

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64} \leftarrow$ 64-bit constant

↑
previous value of rax

x86-64 Semantics

addq $\$0x1$, **%rax**

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb $\$0x1$, **%al**

$\text{al} \leftarrow \text{al} +_8 1_8$

x86-64 Semantics

addq \$0x1, %rax

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb \$0x1, %al

$\text{al} \leftarrow \text{al} +_8 1_8$



x86-64 Semantics

addq \$0x1, %rax

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb \$0x1, %al

$\text{al} \leftarrow \text{al} +_8 1_8$

$\text{rax} \leftarrow \text{rax}[63:8] \circ (\text{rax}[7:0] +_8 1_8)$



x86-64 Semantics

addq \$0x1, %rax

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb \$0x1, %al

$\text{rax} \leftarrow \text{rax}[63:8] \circ (\text{rax}[7:0] +_8 1_8)$

addw \$0x1, %ax

$\text{rax} \leftarrow \text{rax}[63:16] \circ (\text{rax}[15:0] +_{16} 1_{16})$

addl \$0x1, %eax

$\text{rax} \leftarrow \text{rax}[63:32] \circ (\text{rax}[31:0] +_{32} 1_{32})$

x86-64 Semantics

addq \$0x1, %rax

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb \$0x1, %al

$\text{rax} \leftarrow \text{rax}[63:8] \circ (\text{rax}[7:0] +_8 1_8)$

addw \$0x1, %ax

$\text{rax} \leftarrow \text{rax}[63:16] \circ (\text{rax}[15:0] +_{16} 1_{16})$

addl \$0x1, %eax

$\text{rax} \leftarrow 0_{32} \circ (\text{rax}[31:0] +_{32} 1_{32})$

x86-64 Semantics

addq \$0x1, %rax

$\text{rax} \leftarrow \text{rax} +_{64} 1_{64}$

addb \$0x1, %al

$\text{rax} \leftarrow \text{rax}[63:8] \circ (\text{rax}[7:0] +_8 1_8)$

addw \$0x1, %ax

$\text{rax} \leftarrow \text{rax}[63:16] \circ (\text{rax}[15:0] +_{16} 1_{16})$

addl \$0x1, %eax

$\text{rax} \leftarrow 0_{32} \circ (\text{rax}[31:0] +_{32} 1_{32})$

$\text{zf} \leftarrow 0_{32} = (\text{eax} +_{32} 1_{32})$

$\text{cf} \leftarrow ((0_1 \circ \text{eax}) +_{33} 1_{33})[32,32]$

$\text{sf} \leftarrow (\text{eax} +_{32} 1_{32})[31,31]$

$\text{of} \leftarrow \neg \text{eax}[31,31] \wedge (\text{eax} +_{32} 1_{32})[31,31]$

$\text{pf} \leftarrow (\text{eax} +_{32} 1_{32})[0,0] \oplus (\text{eax} +_{32} 1_{32})[1,1] \oplus$
 $(\text{eax} +_{32} 1_{32})[2,2] \oplus (\text{eax} +_{32} 1_{32})[3,3] \oplus$
 $(\text{eax} +_{32} 1_{32})[4,4] \oplus (\text{eax} +_{32} 1_{32})[5,5] \oplus$
 $(\text{eax} +_{32} 1_{32})[6,6] \oplus (\text{eax} +_{32} 1_{32})[7,7]$

Related Work

Manual partial specifications

- CompCert [CACM'09], BAP [CAV'11], BitBlaze [ICISS'08], Codesurfer/x86 [ETAPS'05], McVeto [CAV'10], STOKE [ASPLOS'13], Jakstab [CAV'08], many others

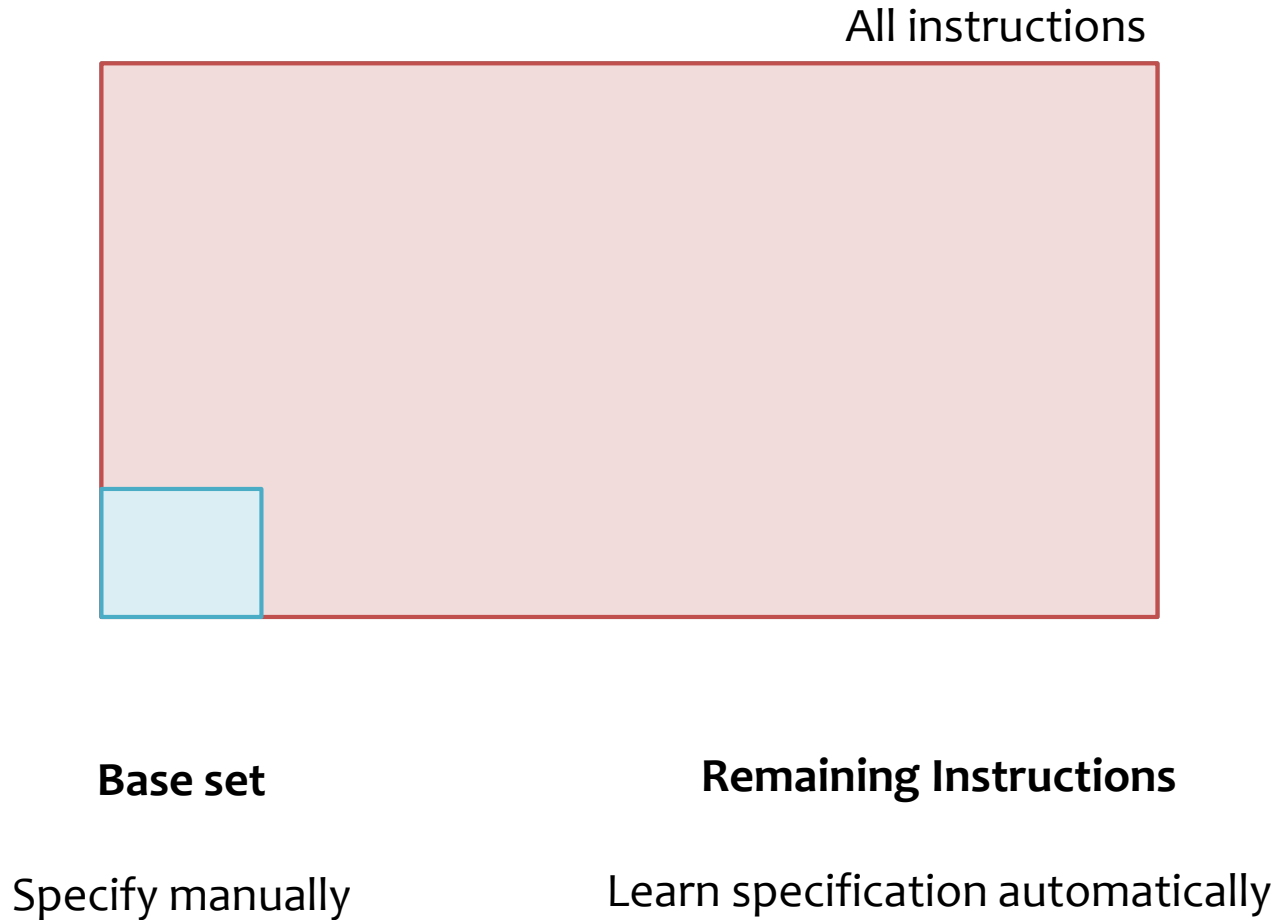
Taly/Godefroid [PLDI'12]

- Automatically synthesize specification from templates
- Only 534 instructions

Automatically Learn a Specification for the x86-64 ISA

Bit-vector formulas of input-output behavior

Strategy: Split Instruction Set



Strategy: Using Program Synthesis



① How do we
synthesize
programs?

② Formal
guarantee?
 $i \equiv \phi$

Strategy: Using Program Synthesis




1

How do we
synthesize
programs?

Randomized search
Guided by cost function
Based on test-cases

Using STOKE [ASPLOS'13]

Strategy: Using Program Synthesis



$$p \equiv \phi$$

2

Formal
guarantee?

$$i \equiv \phi$$

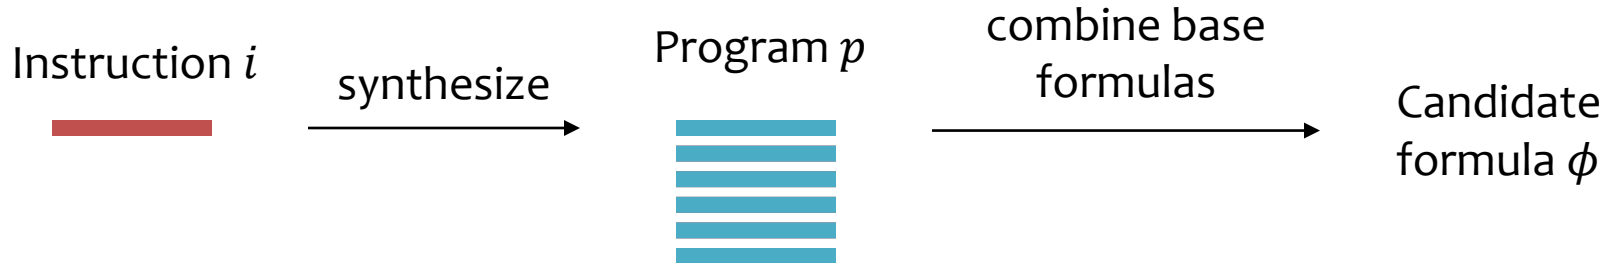
Strategy: Using Program Synthesis



$$i \not\equiv p \equiv \phi$$

2 Formal
guarantee?
 $i \equiv \phi$

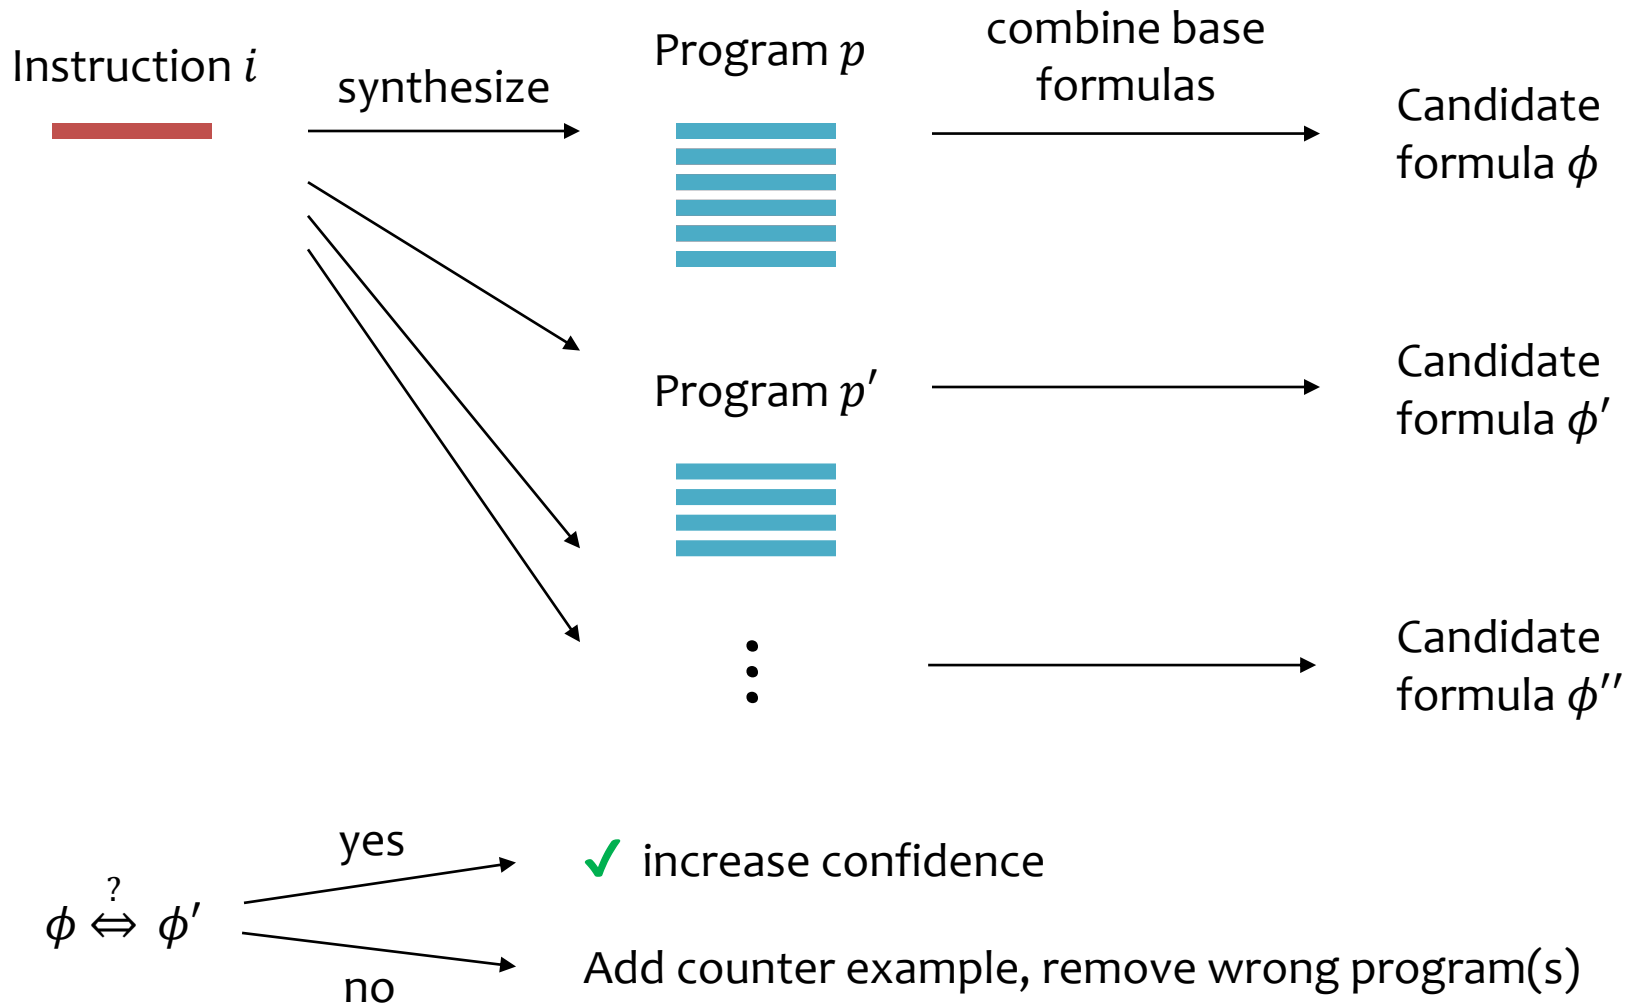
Strategy: Using Program Synthesis



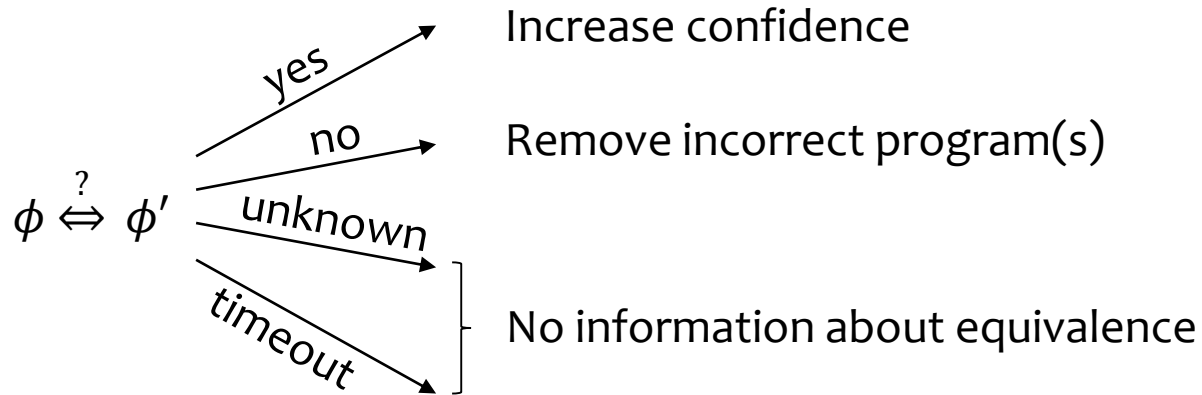
$$i \not\equiv p \equiv \phi$$

2 Formal
guarantee?
 $i \equiv \phi$

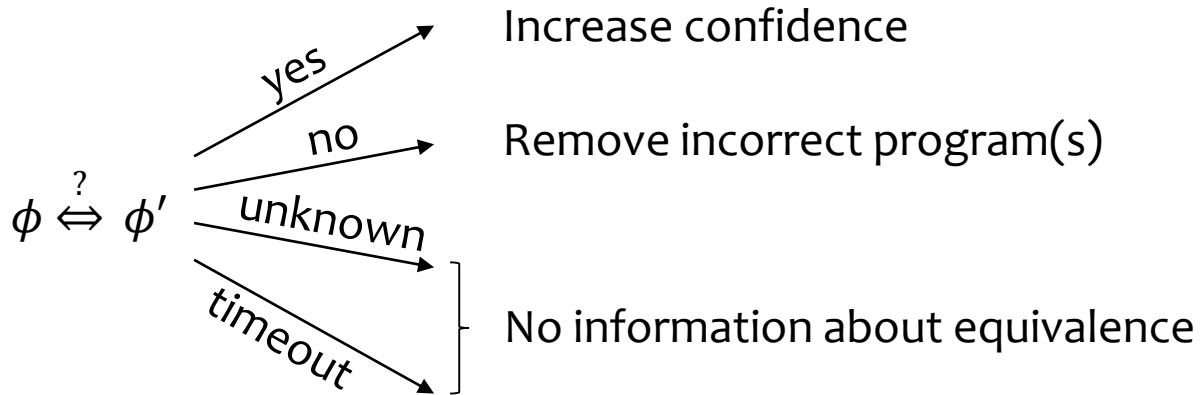
Strategy: Using Program Synthesis




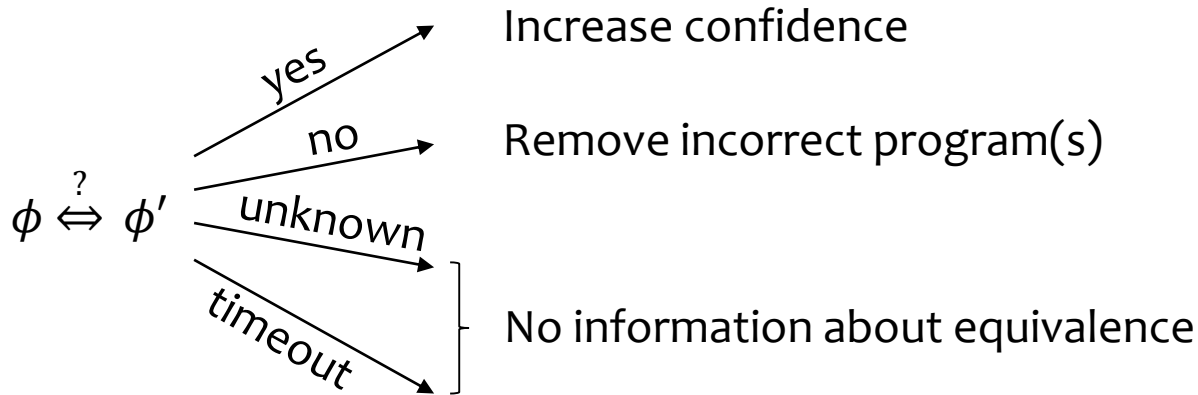
Solver Imprecision



Solver Imprecision



Solver Imprecision



Equivalence class 1



Equivalence class 2

Picking a Program



Equivalence class 1



Equivalence class 2




Equivalence class 3


Prefer programs whose formulas are

- **Precise** (fewest uninterpreted functions)
- **Fast** (fewest non-linear arithmetic operations)
- **Simple** (fewest nodes)

Picking a Program



Equivalence class 1



Equivalence class 2




Equivalence class 3


Prefer programs whose formulas are

- **Precise** (fewest uninterpreted functions)
- **Fast** (fewest non-linear arithmetic operations)
- **Simple** (fewest nodes)

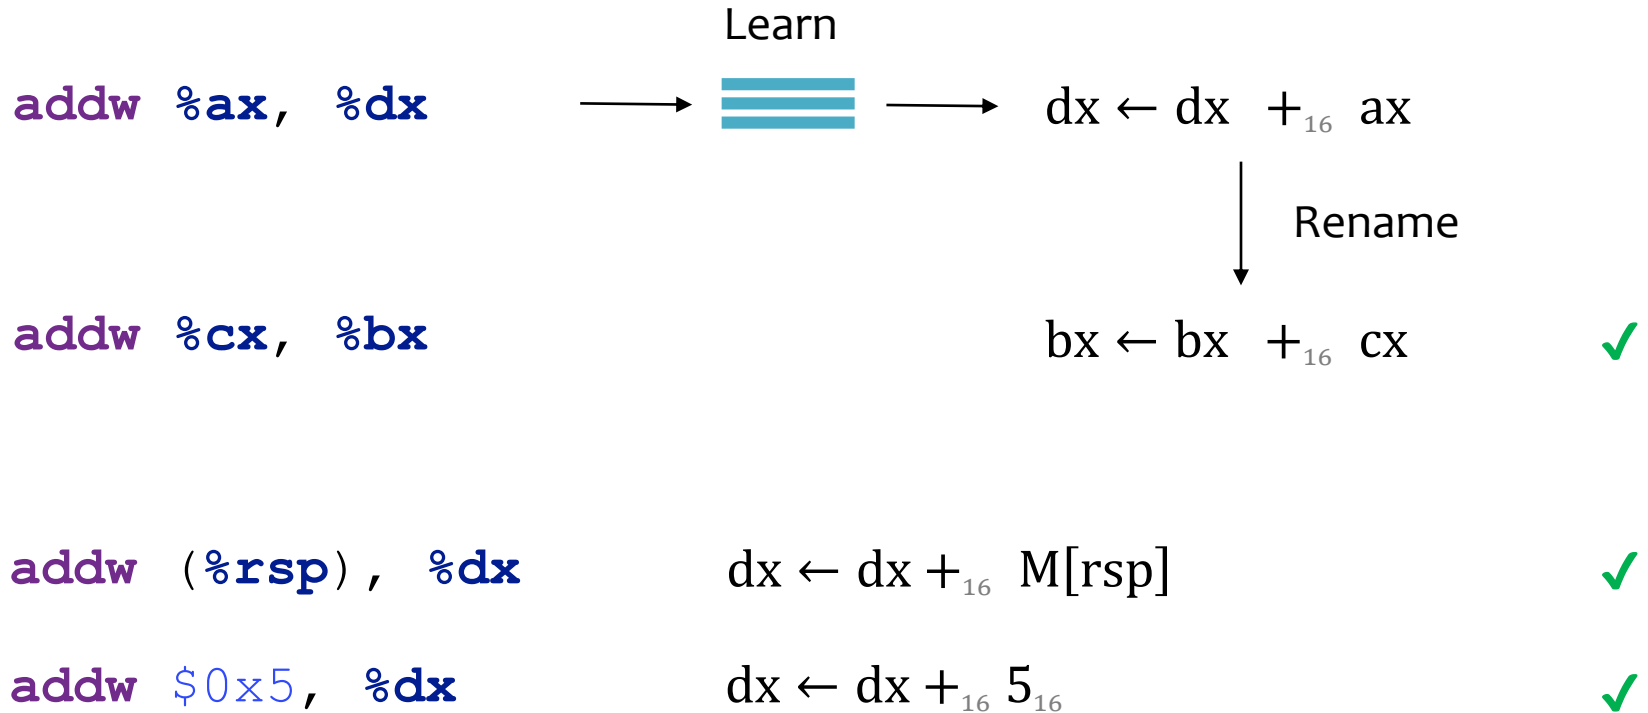
Problem: Synthesis Limitations



Solution: Stratified Search



Generalizing Formulas



Generalization Summary

1. Learn formula for register-only instructions
2. Generalize formulas
 - To other types of operands
3. Check on test inputs

What if Generalization Impossible?

```
shufps $0xb3, %xmm0, %xmm1
```

Problem: No corresponding register-only variant

Solution: Brute force a formula for every constant

Experiment

Base set (51 instructions)

- Integer, bitwise and float operations
- Data movement (including conditional move)
- Conversion operations

Pseudo instructions (11 templates)

- Split and combine registers
- Changing status flags

Goal

Total instructions		3,684
--------------------	--	-------

Out-of-scope

– System instructions	<code>invpcid, jle</code>	302
-----------------------	---------------------------	-----

– Crypto instructions	<code>aeskeygenassist</code>	35
-----------------------	------------------------------	----

– Deprecated instructions	<code>fadd</code>	332
---------------------------	-------------------	-----

– String instructions	<code>scasq</code>	97
-----------------------	--------------------	----

Goal instructions		2,918
-------------------	--	-------

Results

Base set	51
Pseudo instructions	11
<hr/>	
Register-only instructions learned	692
Generalized	984
8-bit constant instructions learned	119.42
<hr/>	
Total formulas learned	1,795.42

Evaluation: Are the Formulas Correct?

Compare with handwritten formulas (from STOKe)

Available for comparison	1,431.91
Automatically proven equivalent	1,377.91
Equivalent with additional lemma	4

Evaluation: Are the Formulas Correct?

Compare with handwritten formulas (from STOKe)

Available

$$\text{fadd}(a, b) = \text{fadd}(b, a)$$

1.91

Automated

7.91

Equivalent

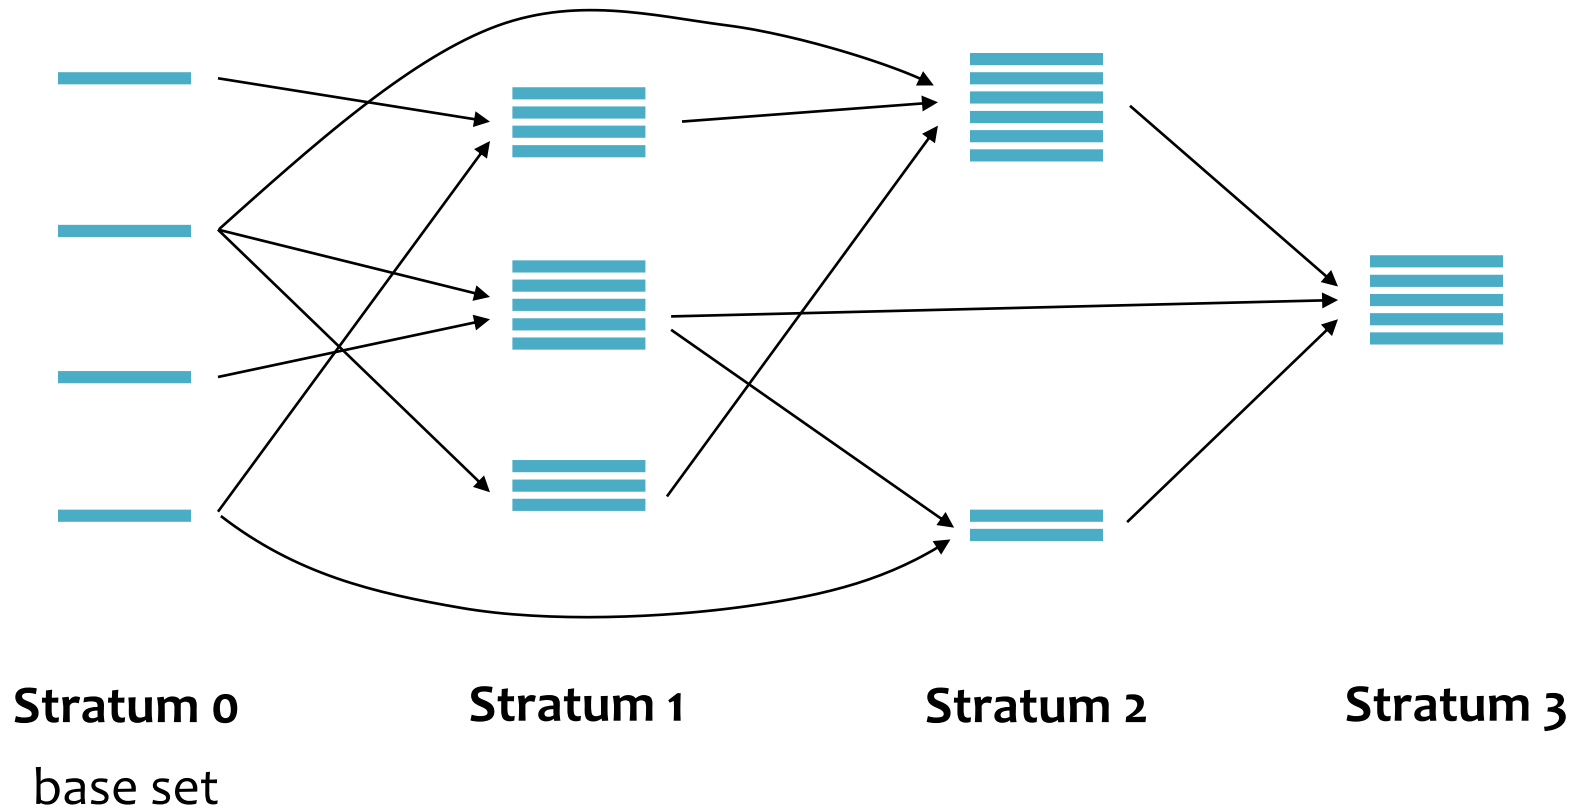
4

Evaluation: Are the Formulas Correct?

Compare with handwritten formulas (from STOKe)


Available for comparison	1,431.91
Automatically proven equivalent	1,377.91
Equivalent with additional lemma	4
Semantically different	50
Handwritten formula correct	0
Learned formula correct	50

Evaluation: Is Stratification Necessary?



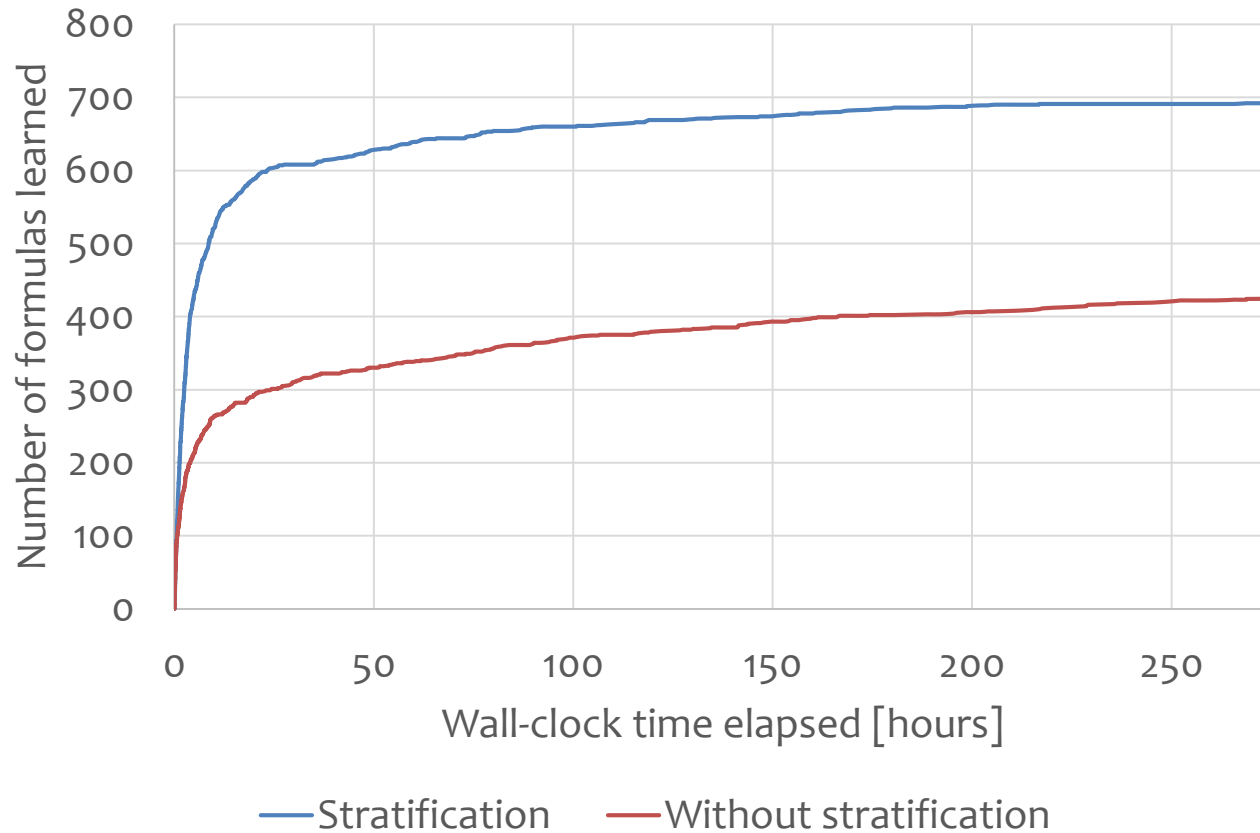
$$\text{stratum}(i) = \begin{cases} 0 & \text{if } i \in \text{baseset} \\ 1 + \max_{i' \in M(i)} \text{stratum}(i') & \text{otherwise} \end{cases}$$

Evaluation: Is Stratification Necessary?

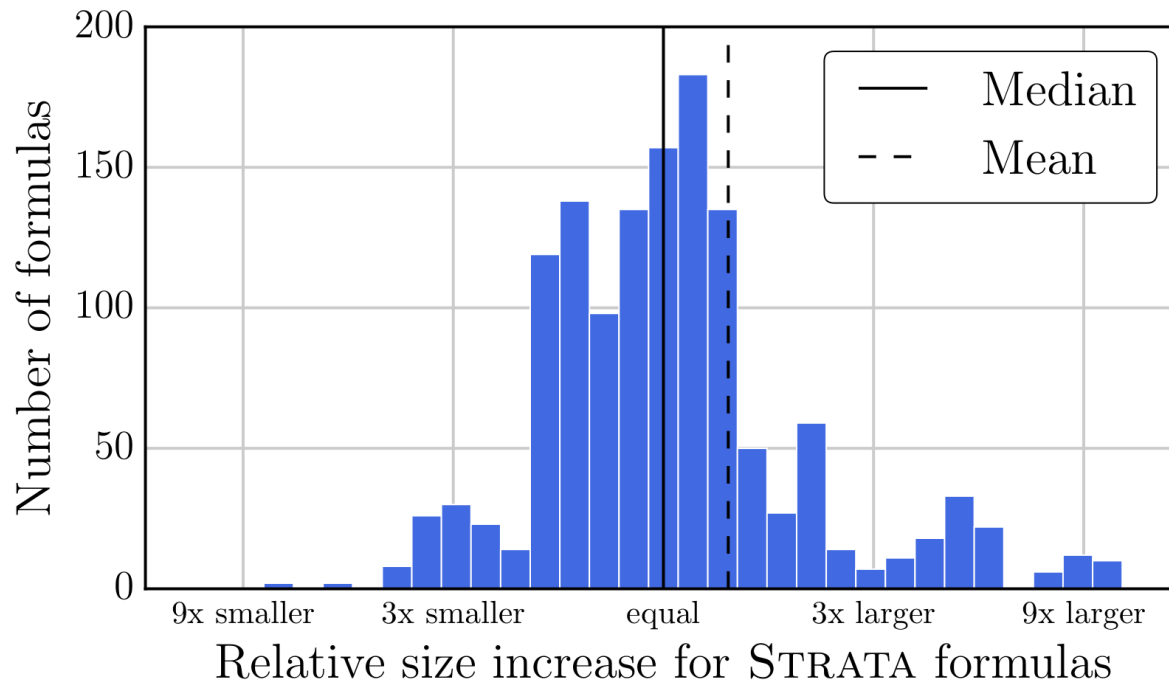
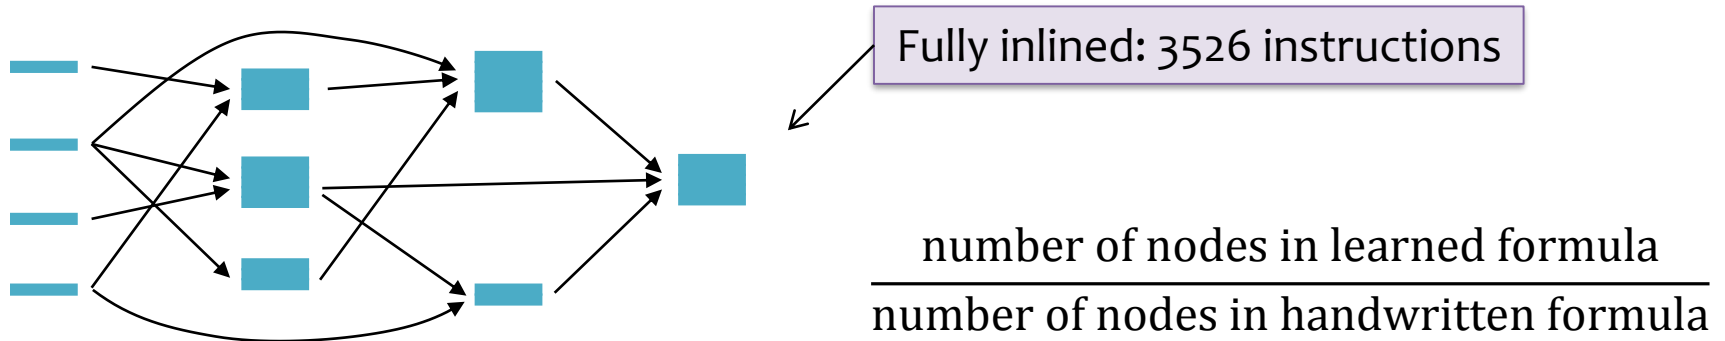


$$\text{stratum}(i) = \begin{cases} 0 & \text{if } i \in \text{baseset} \\ 1 + \max_{i' \in M(i)} \text{stratum}(i') & \text{otherwise} \end{cases}$$

Evaluation: Is Stratification Necessary?



Evaluation: Size of Learned Formulas



Conclusions

1. Automatically learned 1,795 formulas
2. Stratification key to scale program synthesis
3. Compare to hand-written specification
 - More correct, equally precise, same size

Source code, formulas, experimental results



<https://github.com/StanfordPL/strata/>

Backup Slides

Limitations

1. Missing base instructions

Some integer and floating point operations are missing

2. Program synthesis limits

Shortest known program is long and outside of reach
e.g., byte-vectorized operation

3. Cost function limitation

For one bit of output, the cost function does not give enough signal

4. Crazy instructions

SMT solver usage (Z3)

Total decisions	7,075	
Equivalent	6,669	(94.26%)
New equivalence class	356	(5.03%)
Counter-examples	50	(0.71%)
Timeouts (45 seconds):	3	

Experiment Details

Intel Xeon E5-2697 (28 cores) at 2.6 GHz

- 268.86 hours (register-only)
- 159.12 hours (8-bit constants)

Total of 11,983.37 core hours

Test Cases

Random inputs (random machine state)

“Interesting” bit-patterns

0, 1, -1 , 2^n , NaN, Infinity

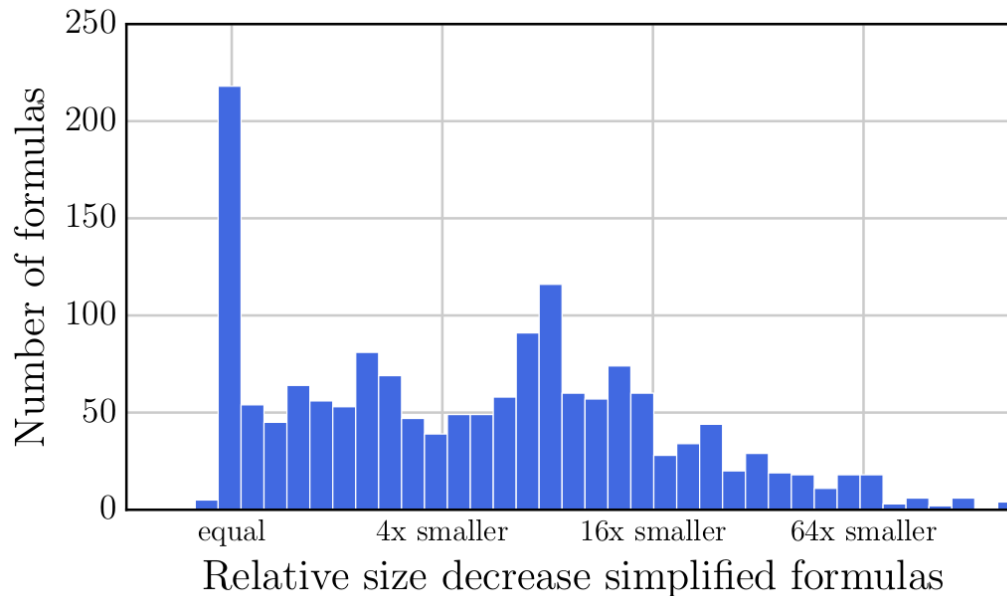
Test cases learned from counter-examples

Evaluation: Simplification

Formulas are simplified

- Constant propagation $2_{64} *_{64} 4_{64} \equiv 8_{64}$
- Move bit-selection over concatenation

$$(0_{64} \circ \text{rax})[63,0] \equiv \text{rax}$$



Evaluation: Formula Precision

Formula precision (number of uninterpreted functions)

- Learned formulas equally precise in all but 4 cases

Formula quality (number of non-linear operations)

- Learned formulas contain same number of non-linear operations, except for 11 cases